

Databases Concepts

Objectives

- Why RDBMS?
- Different models (hierarchy, network, OO, relational)
- RDBMS concepts: tables, keys, relationships
- ER model concepts: entities, attributes, relationships, cardinality
- Mapping entity relationships to tables
- Normalization
- Using MySQL
- Data types in MySQL
- Create tables, basic DDL
- DDL commands (tables, relationships, keys, constraints)

Objectives

- DML (Insert, update, delete)
- Queries (select, predicates, identity)
- Complex queries
- Built-in MySQL functions
- Grouping and ordering
- Joins
- Subqueries

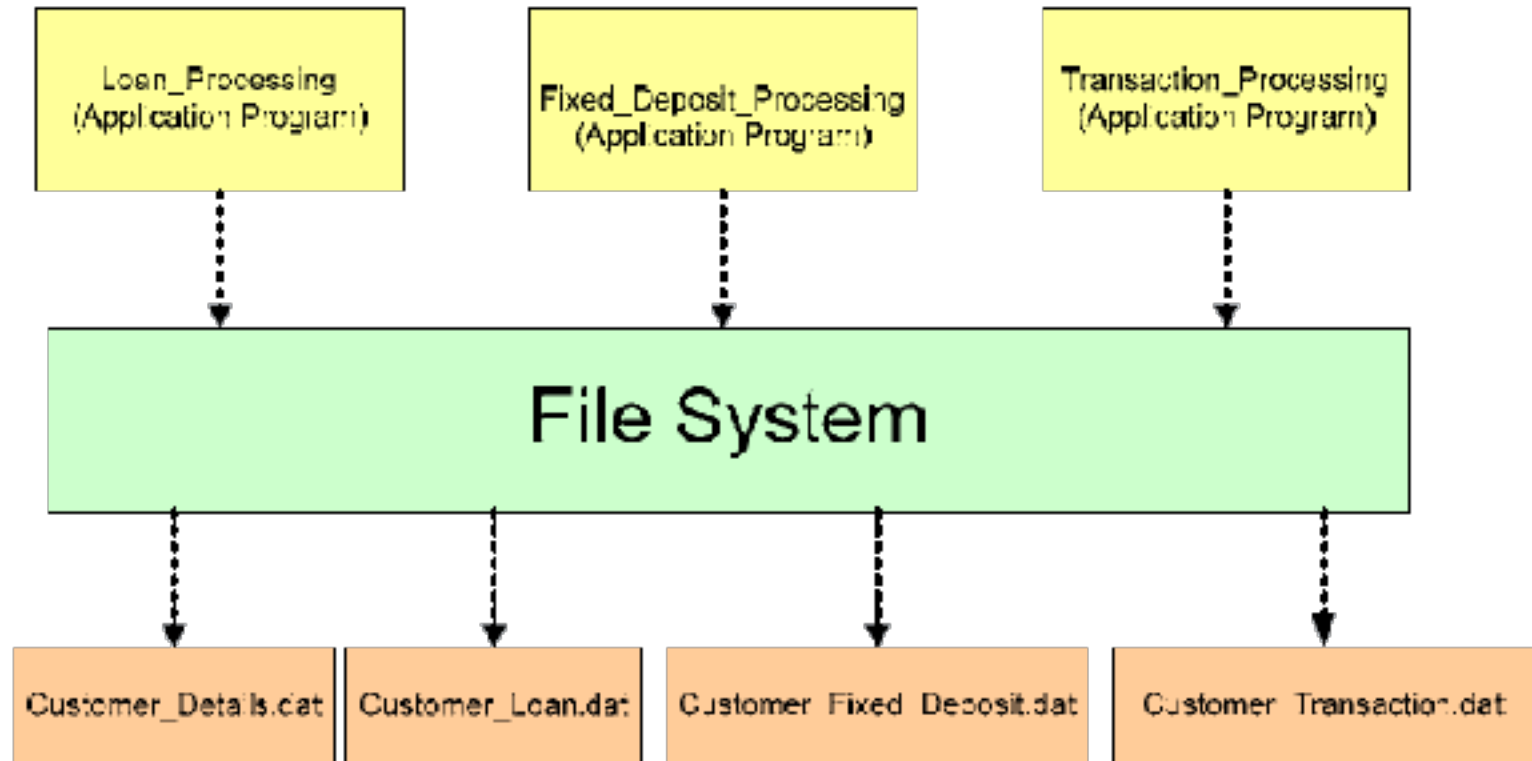
What is

- Data?
 - Meaningful facts, text, graphics, images, sound, video segments.
 - Database?
 - An organized collection of logically related data.
 - Information?
 - Data processed to be useful in decision making.
 - Metadata?
 - Data that describes data.
-

Database

- Represents some aspect of the real world
 - Also known as the miniworld
- Changes to miniworld are reflected in the database
- Database has some source
 - some degree of interaction with events in real world

Traditional method of data storage



Limitations of traditional approach

- Data Dependence
 - Data Redundancy (Duplication of data)
 - Limited Data Sharing
 - Lengthy Development Times
 - Excessive Program Maintenance
-

Why DBMS?

- Controlling redundancy
 - Sharing of data
 - Restricting unauthorized access
 - Providing multiple interfaces
 - Enforcing integrity constraints
 - Providing backup & recovery
 - Potential for enforcing standards
 - Reduced application development time
 - Availability of up-to-date Information
-

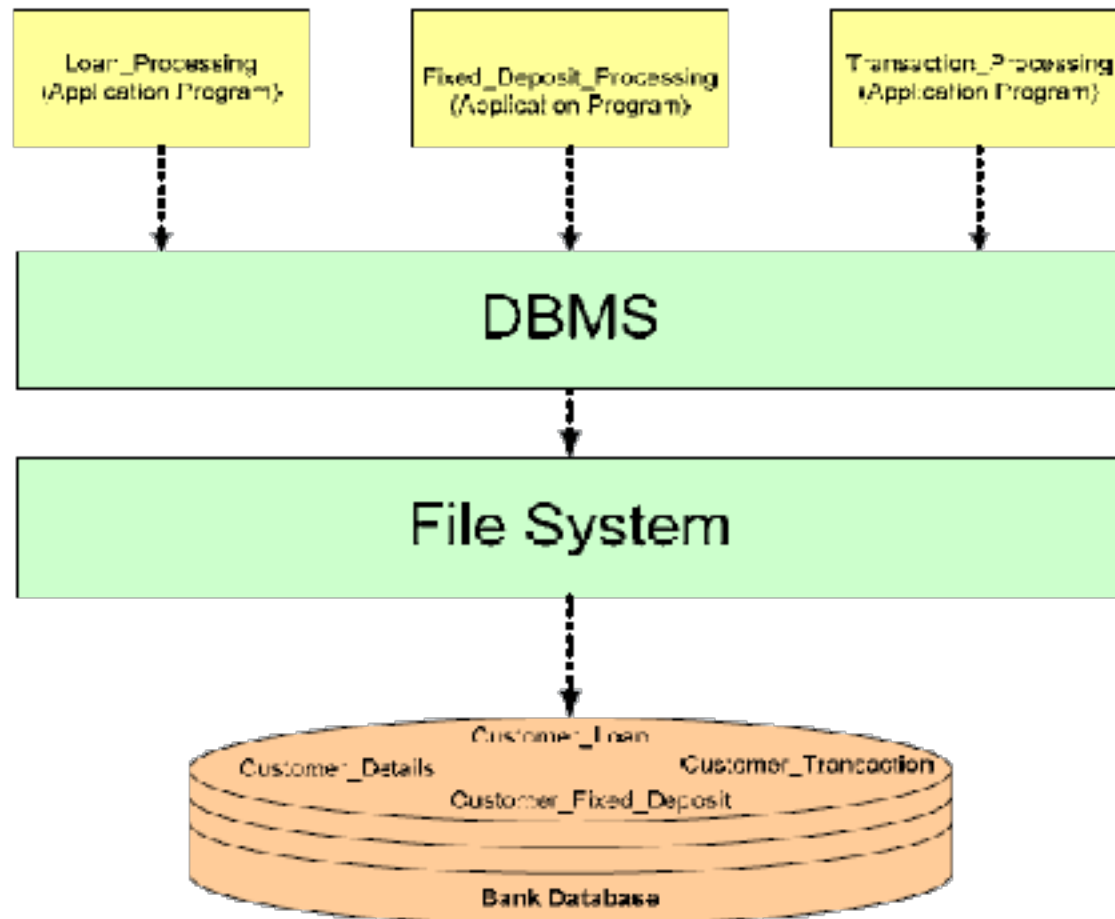
Database Management System

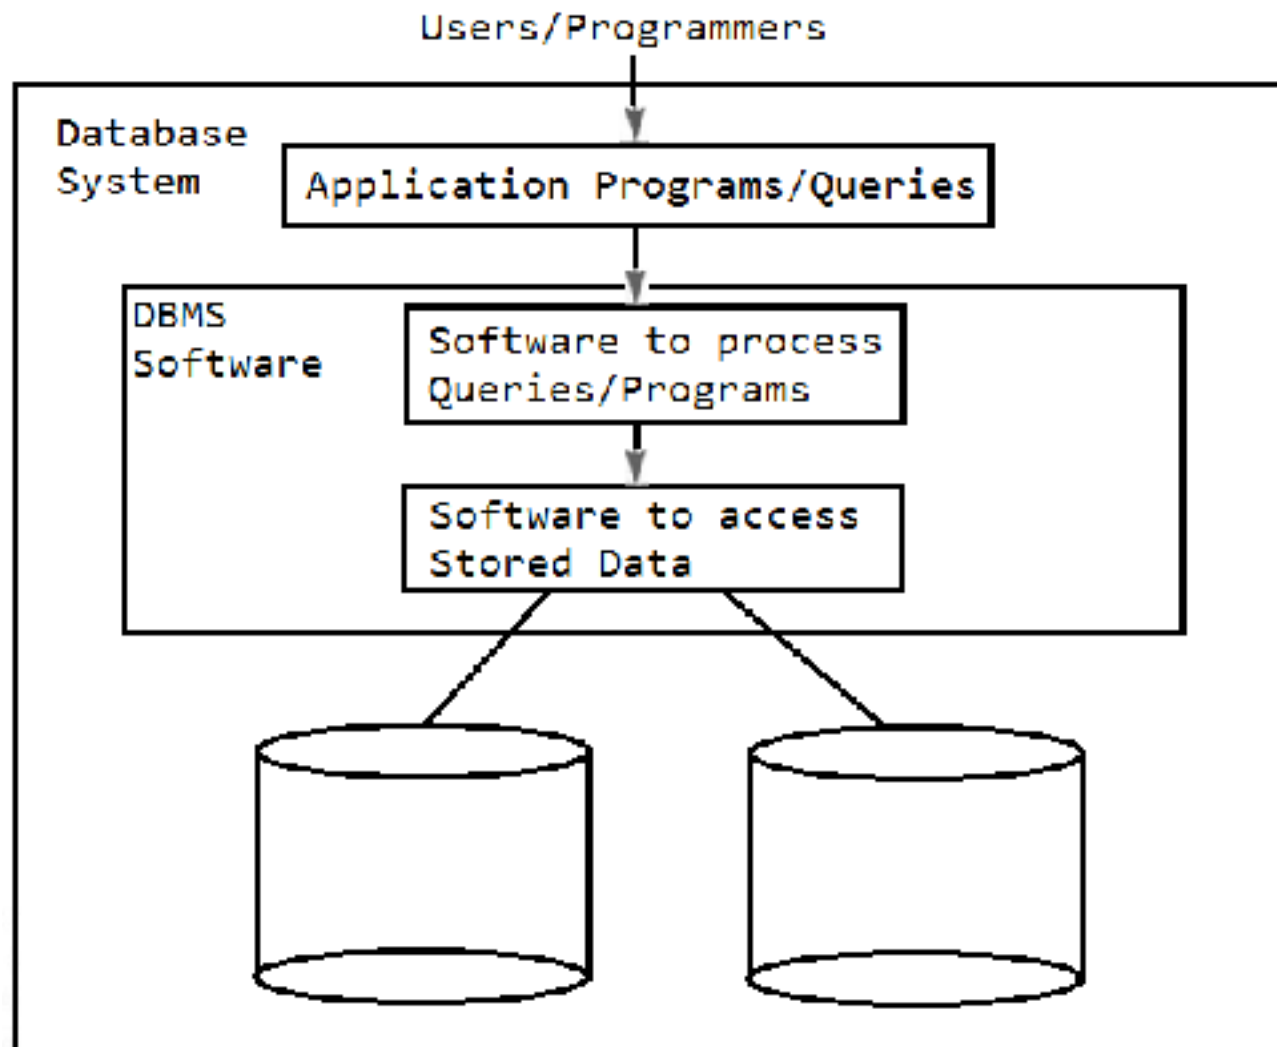
- Database Management System
 - A computerized record-keeping system
 - A data storage and retrieval system which permits data to be stored non-redundantly while making it appear to the user as if the data is well-integrated.
 - General purpose software system
 - Facilitates the processes of
 - Defining, constructing, and manipulating databases for various applications
-

Database Management System

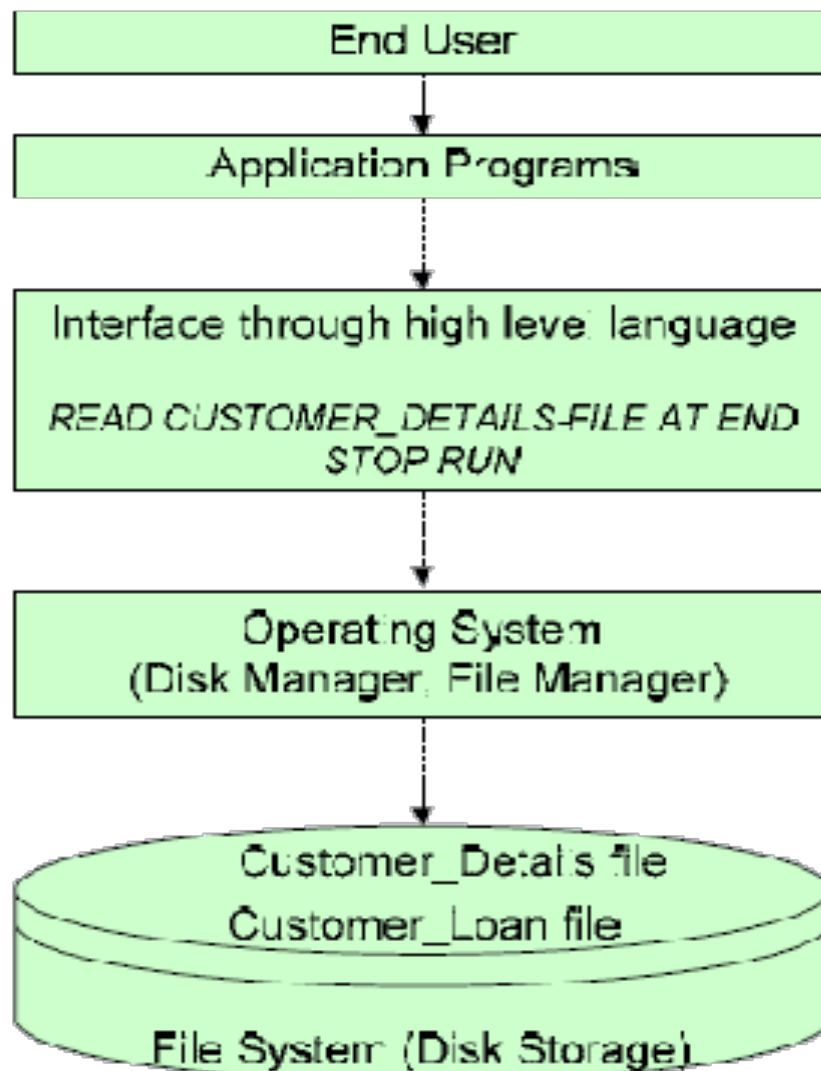
- Defining
 - specifying types of data and detailed description of each type
 - Constructing
 - storing the data itself on some storage medium
 - Manipulating:
 - Querying
 - Updating
 - Generating reports
-

Position of DBMS

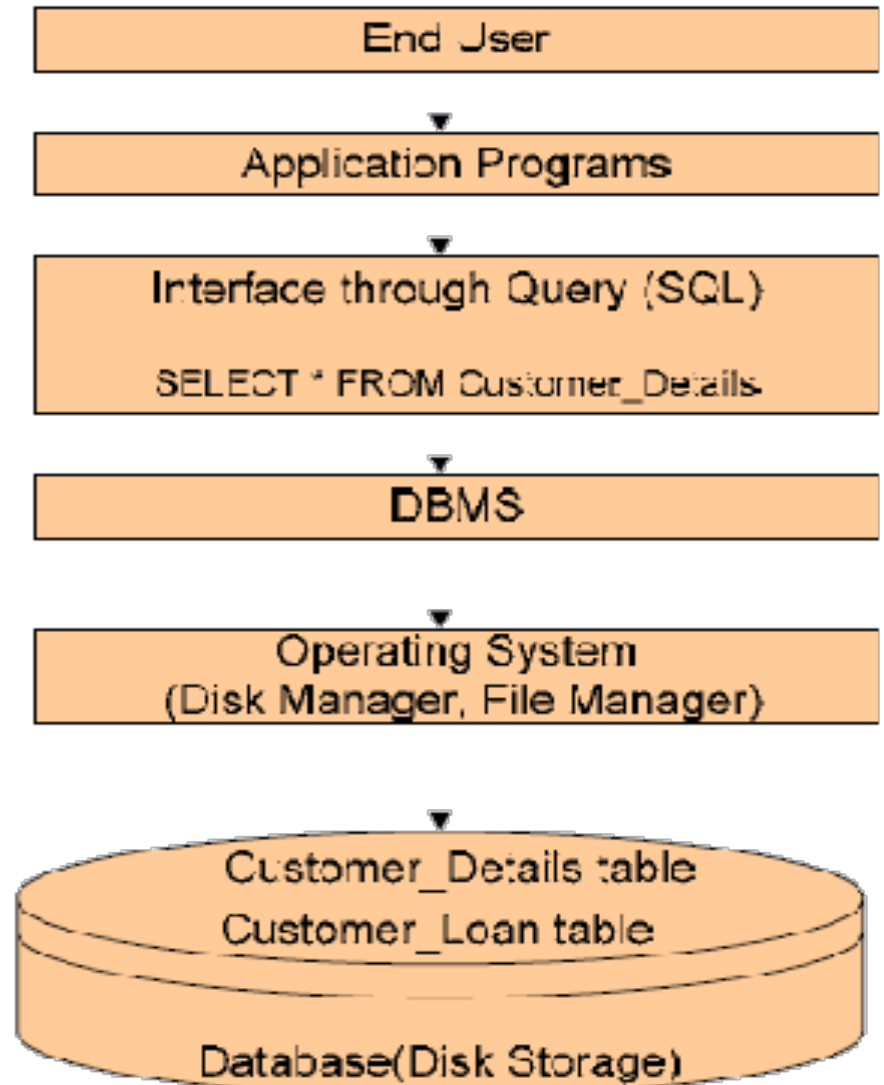


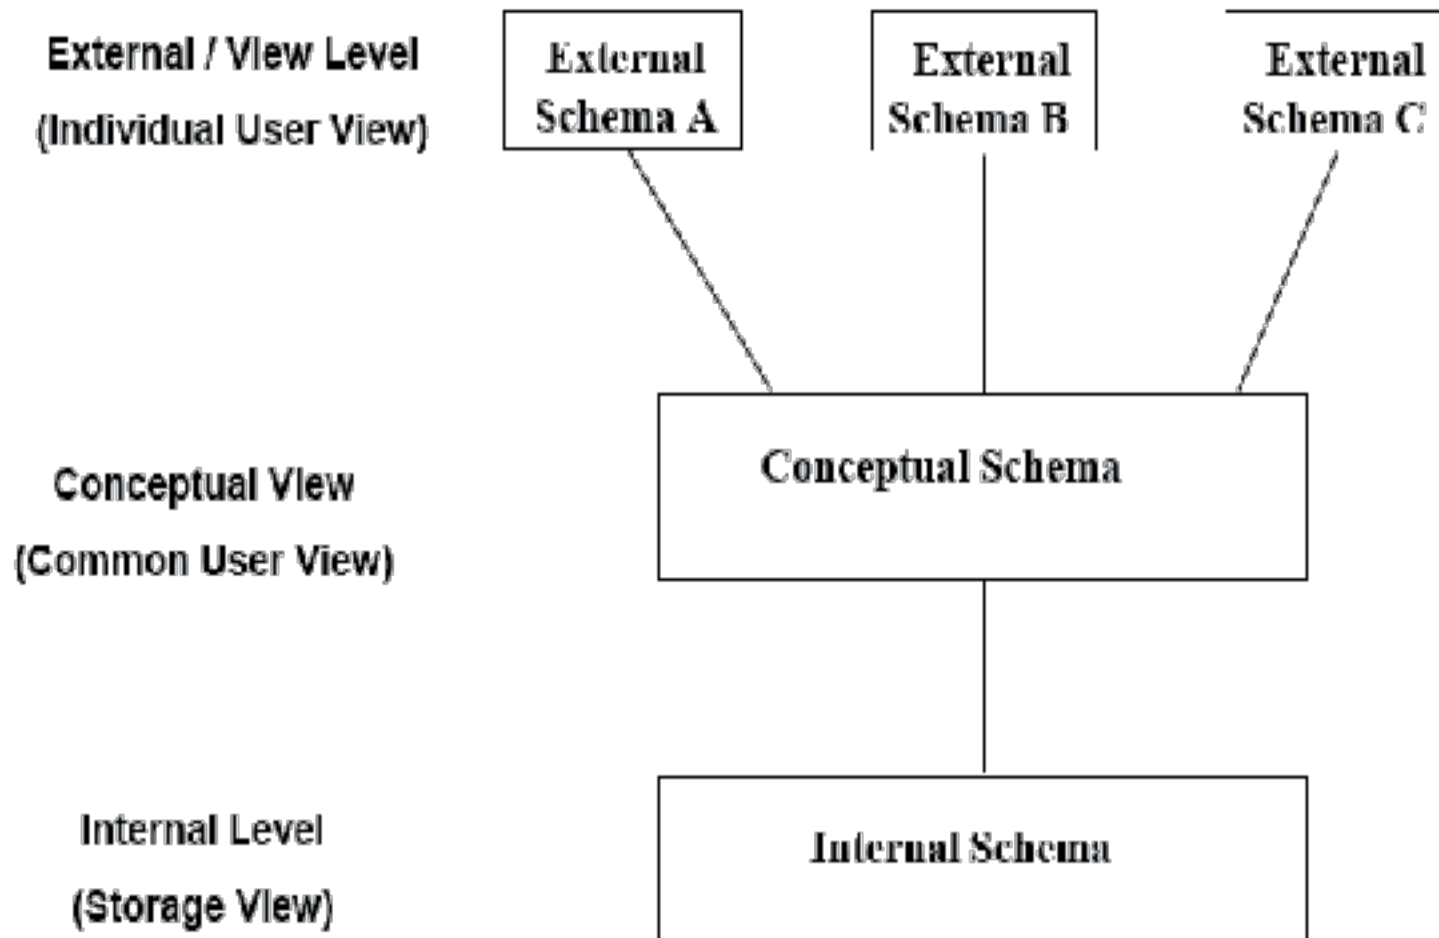


File system Interface



DBMS Interface





Customer_Loan Cust_ID : 101 Loan_No : 1011 Amount in Dollars : 8755.00	<i>External</i>
CREATE TABLE Customer_Loan (Cust_ID NUMBER Loan_No NUMBER(4) Amount_in_Dollars NUMBER(7,2))	<i>Conceptual</i>
Cust ID TYPE = BYTE (4), OFFSET = 0 Loan_No TYPE = BYTE (4), OFFSET = 4 Amount in Dollars TYPE = BYTE (7), OFFSET = 8	<i>Internal</i>

Data Model (Self Learning)

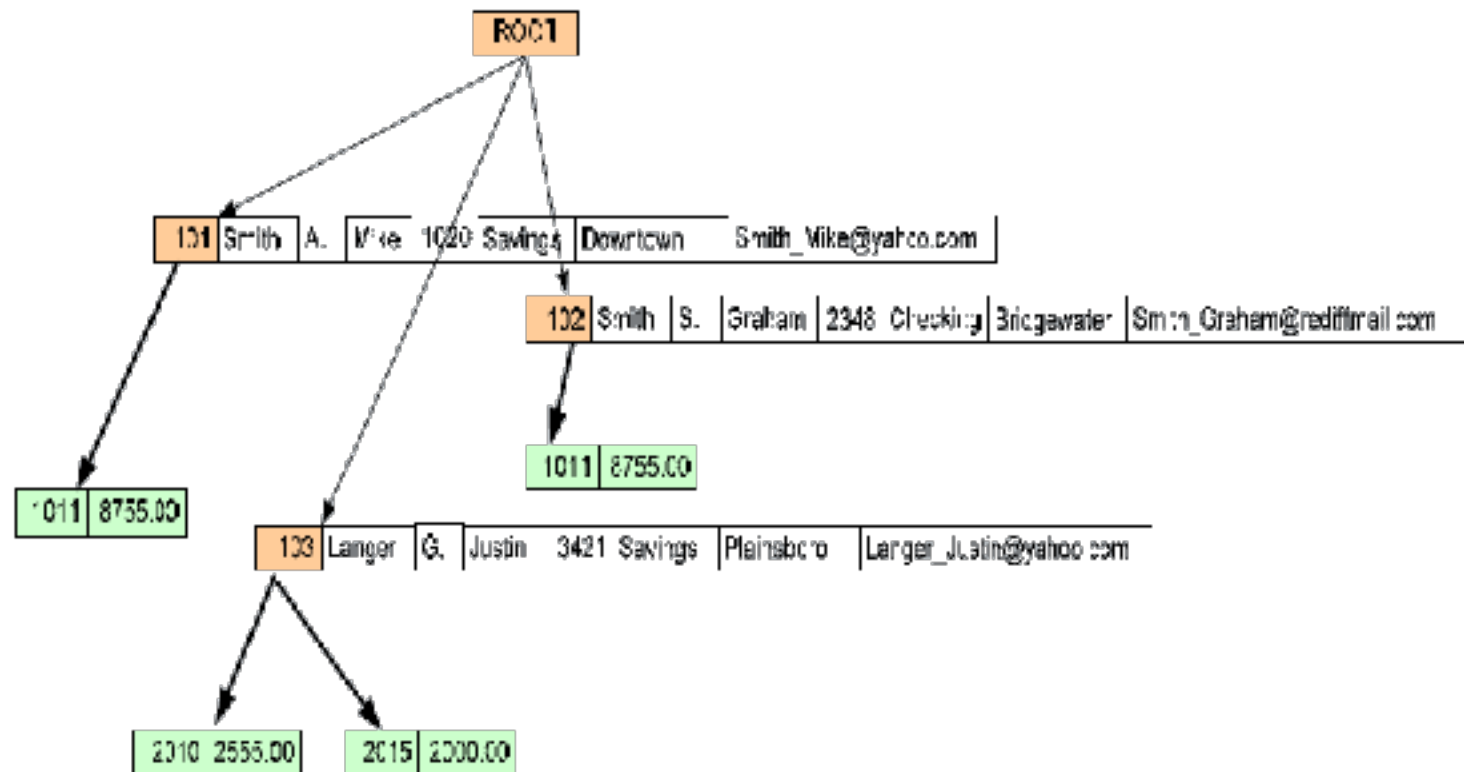
- Is a collection of concepts that can be used to describe the structure of a database
 - High-level/Conceptual Model-provides concepts that are close to the way many users perceive data
 - Low-level/physical Models-provides concepts that describe the details of how data is stored on computer
 - Representational/implementation Models-provides concepts understood by end users , hides few details on how data is stored
-

Evolution of Database Systems (Implementation Model)

- Hierarchical and Network (Legacy)
 - Relational
-

(Self Learning)

Record based data model – Hierarchical data model



Drawbacks of Hierarchical Model (Self Learning)

- Knowledge of the structure and its physical representation is required.
 - Does not support for modifying the structure
 - Time consuming as the processing is record at a time.
-

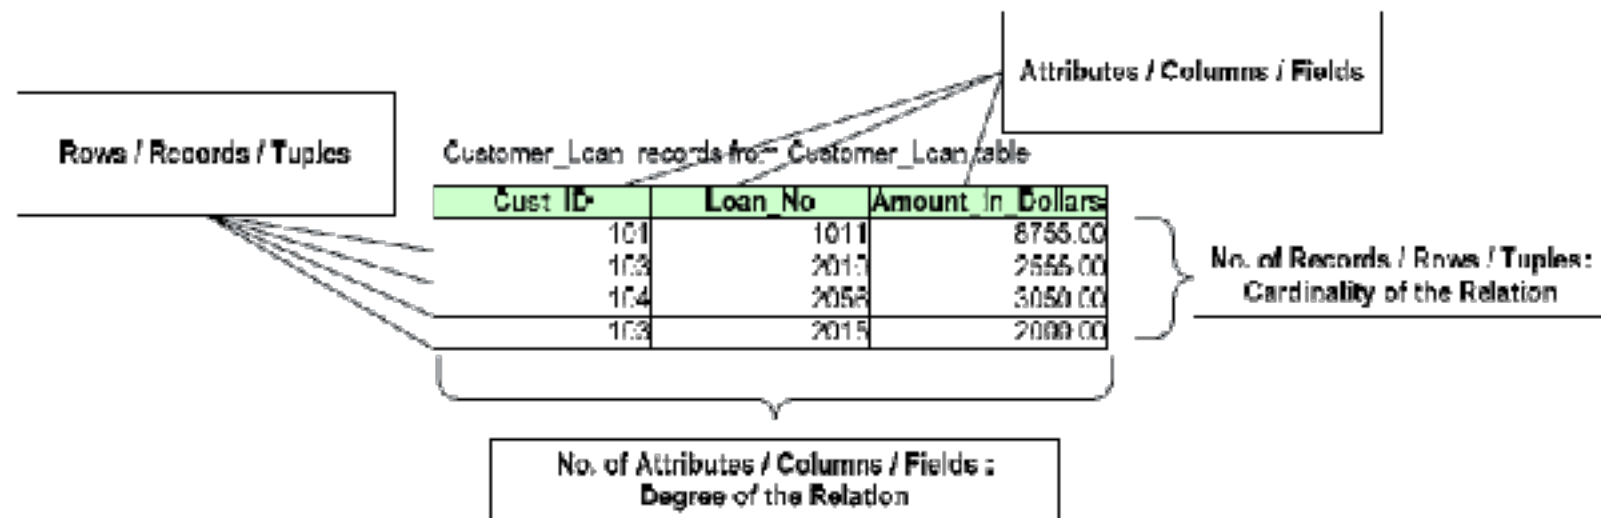
Record based data model – Network data model



Drawbacks of Network Model (Self Learning)

- Structure is inherently complex.
 - Knowledge of the structure and its physical representation is required.
 - Does not support for modifying the structure
 - time consuming as the processing is record at a time.
-

Record based data model – Relational data model



Cust_ID	Cust_Last Name	Cust_Mid _Name	Cust_First _Name	Account _No	Account Type	Bank_Branch	Cust_Email
101	Smith	A.	Mike	1020	Savings	Downtown	Smith_Mike@yahoo.com
102	Smith	S.	Graham	2345	Checking	Bridgewater	Smith_Graham@rediffmail.com
103	Langer	C.	Justin	3421	Savings	Plainsboro	Langer_Justin@yahoo.com
104	Quails	D.	Jack	2387	Checking	Downtown	Quails_Jack@yahoo.com
105	Jones	E.	Simon	2309	Checking	Brighton	Jones_Simon@rediffmail.com

Customer_Detail records from Customer_Details table

RDBMS concepts: tables, keys, relationships

- Relation, Tuple, Attribute
 - Every relation has a unique name.
 - Every attribute value is atomic.
 - Every row/tuple is unique.
 - Attributes in tables have unique names.
 - Sometimes the value for a particular attribute may be unknown, or it may have no value. This is represented by a null
 - Null is not the same as zero, blank or an empty string
-

RDBMS concepts: tables, keys, relationships

- Relational Keys and Structures
 - Keys
 - Candidate Key is a set of one or more attributes that can uniquely identify a row in a given table.
 - Primary Key is an attribute of an entity which can uniquely identify a row in a given table
 - Composite Key is a primary key which is a combination of more than one attribute is called a composite primary key
 - Foreign Key is an attribute of an entity which refers to the domain of the primary key in the same entity or a different entity
-

Integrity Constraints

- Domain Integrity
 - Allowable values for an attribute. (not null, check)
 - Entity Integrity
 - No primary key attribute may be null. (unique, not null)
 - Referential Integrity
 - For example: Delete Rules
 - Restrict, Cascade, Set-to-Null
 - To be discussed more in SQL
-

Schema for Relations

CUSTOMER

<u>Customer_ID</u>	Customer_Name	Address	City	State	Zip
--------------------	---------------	---------	------	-------	-----

ORDER

<u>Order_ID</u>	Order_Date	<u>Customer_ID</u>
-----------------	------------	--------------------

ORDER_LINE

<u>Order_ID</u>	<u>Product_ID</u>	Quantity
-----------------	-------------------	----------

PRODUCT

<u>Product_ID</u>	Product_Description	Product_Finish	Unit_Price	On_Hand
-------------------	---------------------	----------------	------------	---------

Conceptual Data Modeling

The Entity-Relationship Model

E-R Model Constituents

- Entity
 - Person, Place, Object, Event, Concept...
- Attributes
 - Name, Age, City, Salary...
- Relationships
 - Student registers for a Course
 - Employee applies for a Leave

Simple Vs composite attribute

- Simple Attribute:
 - Can not be divided into simpler components
 - For example, age of a Person
 - Composite Attribute:
 - Can be split into components
 - Address of a Person can be split into street, location, city, state, zip
 - Date of birth of an Employee can be split into day, month, and year
-

Single Vs Multi-values Attributes

- Single Valued:
 - Can take only a single value for each entity instance
 - For example, Salary of an employee (only one value)
 - Multi-valued:
 - Can take many values
 - For example, Skill-Set of an employee can be “Java, Servlet, JSP” or “C#, ASP.Net, WCF, WPF”
-

Stored Vs Derived Attributes

- Stored Attribute:
 - Attribute that need to be stored permanently
 - For example, name of an employee
 - Derived Attribute:
 - Attribute that can be calculated based on other attributes
 - For example, years-of-service of an employee can be calculated from the date of hire and the current date
-

Regular Vs Weak Entity

- Regular Entity:
 - Entity that has its own key attribute
 - For example, Employee, Student, Customer etc.
 - Weak Entity:
 - Entity that depends on other entity for its existence and does not have key attribute of its own
 - For example, Nominee of an Insurance policy holder.
-

Relationships

- A relationship type between two entity types defines the set of all associations between these entity types
 - Each instance of the relationship between the members of these entity types is called a relationship instance
-

Degree of a Relationship

- Degree: The number of entity types involved
 - Unary One
 - Binary Two
 - Ternary Three
 - For example,
 - Employee and Manager-of-Employee is Unary
 - Employee works for Department is Binary
 - Customer purchases Item from a Shop is Ternary
-

Cardinality

- Relationships can have different connectivity
 - One-To-One (1:1)
 - One-To-Many (1:N)
 - Many-To-One (N:1)
 - Many-To-Many (N:M)
-

Cardinality

- One-To-One
 - One employee is given only one laptop and one laptop has only one owner
 - One department has only one head-of-department and each HOD heads only one department
-

Cardinality

- One-To-Many
 - One department has many employees
 - Many-To-One
 - Many employees belong to one department
-

Cardinality

- Many-To-Many
 - One project has many employees, and one employee has worked in many projects
 - One customer has purchased many products, and one product is purchased by many customers
-

Relationship Participation

- Total: Every entity instance must be connected through the relationship to another instance of the other participating entity types
- Partial: All instance need not participate
 - For example,
Employee and Head-Of-Department
Employee – Partial
Department - Total

All employees will not be head-of some department. So only few instances of employee entity participate in the above relationship. But each department will be headed by some employee. So department entity's participation is total and employee entity's participation is partial in the above relationship



Entity

An Entity is an object or concept about which business user wants to store information.



Entity

A weak Entity is dependent on another Entity to exist. Example Order Item depends upon Order Number for its existence. Without Order Number it is impossible to identify Order Item uniquely.



Attribute

Attributes are the properties or characteristics of an Entity



Attribute

A key attribute is the unique, distinguishing characteristic of the Entity



Attribute


A multivalued attribute can have more than one value. For example, an employee Entity can have multiple skill values.




A derived attribute is based on another attribute. For example, an employee's monthly salary is based on the employee's basic salary and House rent allowance.



Relationships illustrate how two entities share information in the database structure.



Attributes



**Multi-valued
Attribute**



Key Attribute



Composite Attribute



Relationship




Unary Relationship



Binary Relationship



Ternary Relationship



Relationship Participation

Weak Entity




Attributes Of Relationship

Transforming E-R Diagrams into Relations

- Map Regular Entities to Relations
- Composite attributes: Use only their simple, component attributes
- Multi-valued Attribute - Becomes a separate relation with a foreign key taken from the superior entity
- Map Weak Entities
 - Becomes a separate relation with a foreign key taken from the superior entity

Map Binary Relationships

Binary 1:1



- Case 1: Combination of participation types

The primary key of the partial participant will become the foreign key of the total participant

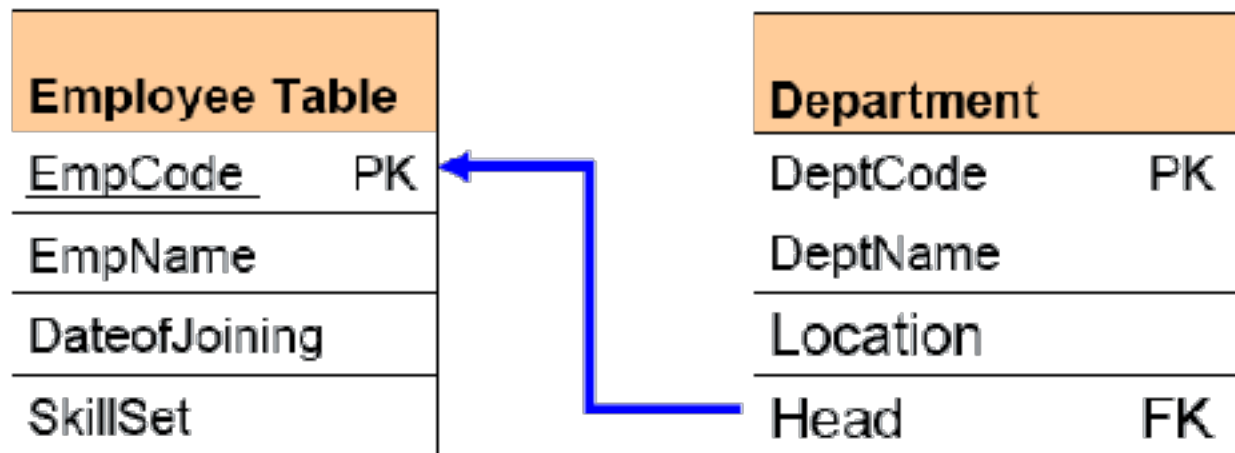
Employee(E#, Name,...)

Department(Dept#, Name..., Head)



Binary 1:1

Binary 1 : 1



Map Binary Relationships

Binary 1:N



The primary key of the relation on the "1" side of the relationship becomes a foreign key in the relation on the "N" side

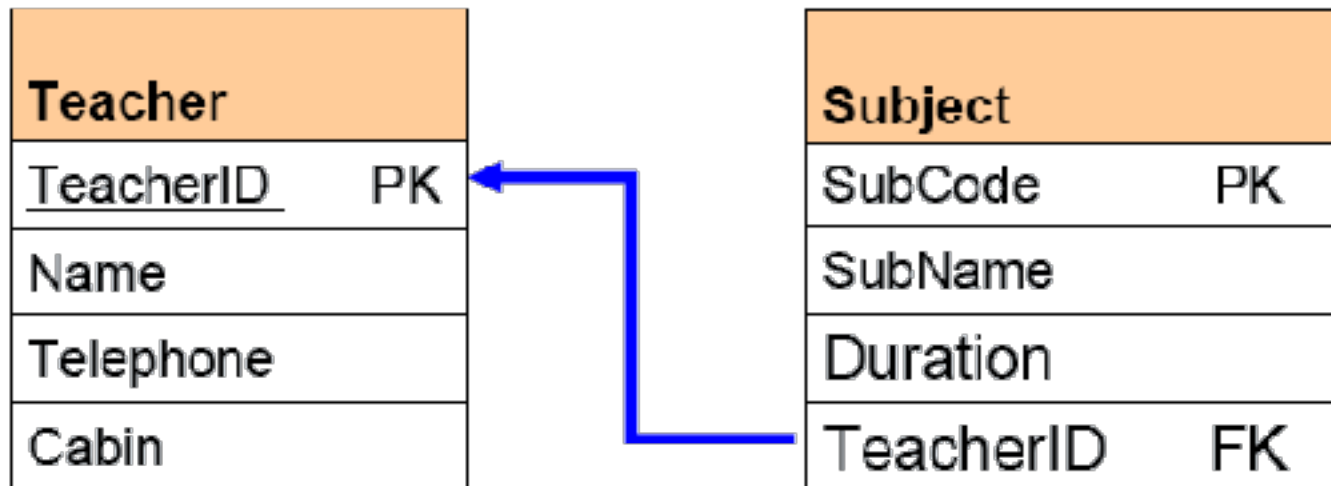
Teacher (ID, Name, Telephone, ...)

Subject (Code, Name, ..., Teacher)




Binary 1:N

Binary 1 : N



Binary M:N

Binary M:N



- A new table is created to represent the relationship
- Contains two foreign keys - one from each of the participants in the relationship
- The primary key of the new table is the combination of the two foreign keys

Student (Sid#, Title...)

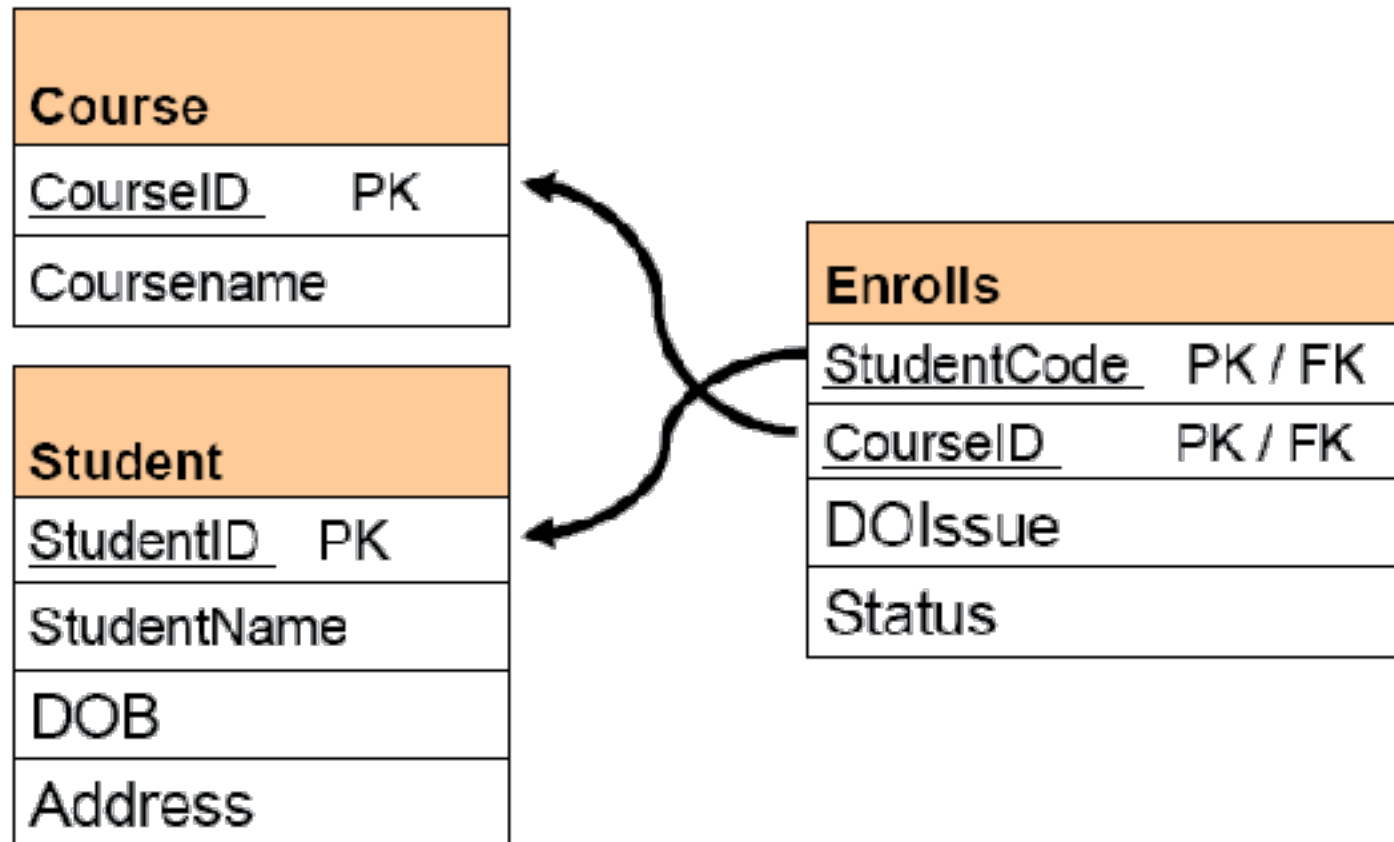
Course(C#, CName,...)

Enrolls (Sid#, C#)



Binary M:N

Binary M : N




Unary

Unary 1:1

- Consider employees who are also a couple
- The primary key field itself will become foreign key in the same table

Employee(E#, Name,... Spouse)



Unary 1 : N

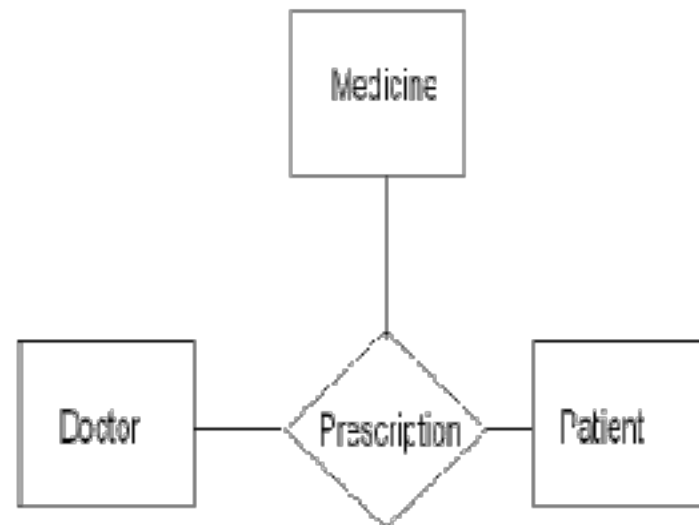
Employee Table	
<u>EmpCode</u>	PK
EmpName	
DateofJoining	
SkillSet	
Manager	FK



(Self Learning)

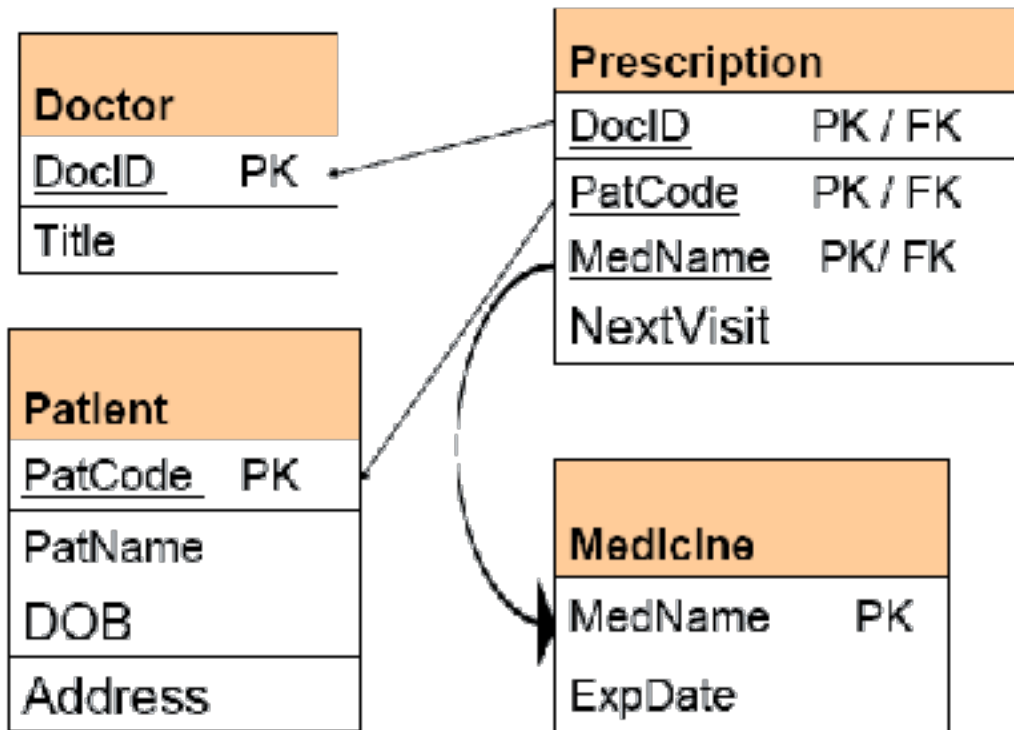
Ternary relationship

- Represented by a new table
- The new table contains three foreign keys - one from each of the participating Entities
- The primary key of the new table is the combination of all three foreign keys
- Prescription (Doctor#, Patient #, Medicine Name)



(Self Learning)

Ternary



Redundancy and normalization

- The Evils of Redundancy
 - Redundancy is at the root of several problems associated with relational schemas:
 - redundant storage, insert/delete/update anomalies
 - Integrity constraints, in particular functional dependencies, can be used to identify schemas with such problems and to suggest refinements.
 - Main refinement technique: decomposition
 - replacing ABCD with, say, AB and BCD, or ACD and ABD.
 - Decomposition should be used judiciously:
 - Is there reason to decompose a relation?
 - What problems (if any) does the decomposition cause?
-

Redundancy and normalization

What is Normalization?

- Database designed based on the E-R model may have some amount of
 - Inconsistency
 - Uncertainty
 - Redundancy

To eliminate these draw backs some **refinement** has to be done on the database.

- **Refinement** process is called **Normalization**
- Defined as a step-by-step process of decomposing a complex relation into a simple and stable data structure.
- The formal process that can be followed to achieve a good database design
- Also used to check that an existing design is of good quality
- The different stages of normalization are known as “normal forms”
- To accomplish normalization we need to understand the concept of Functional Dependencies.

Redundancy and normalization

Student_Course_Result Table

Student_Details			Course_Details				Result_Details		
101	Davis	11/4/1986	M4	Applied Mathematics	Basic Mathematics	7	11/11/2004	82	A
102	Daniel	11/6/1987	M4	Applied Mathematics	Basic Mathematics	7	11/11/2004	62	C
101	Davis	11/4/1986	H6	American History		4	11/22/2004	79	B
103	Sandra	10/2/1988	C3	Bio Chemistry	Basic Chemistry	11	11/16/2004	65	B
104	Evelyn	2/22/1986	B3	Botany		8	11/26/2004	77	B
102	Daniel	11/6/1987	P3	Nuclear Physics	Basic Physics	13	11/12/2004	68	B
105	Susan	8/31/1985	P3	Nuclear Physics	Basic Physics	13	11/12/2004	89	A
103	Sandra	10/2/1988	B4	Zoology		5	11/27/2004	54	D
105	Susan	8/31/1985	H6	American History		4	11/22/2004	87	A
104	Evelyn	2/22/1986	M4	Applied Mathematics	Basic Mathematics	7	11/11/2004	65	B

Functional Dependencies and Keys

- In a given relation R , X and Y are attributes. Attribute Y is **functionally dependent** on attribute X if each value of X determines **EXACTLY ONE** value of Y , which is represented as $X \rightarrow Y$ (X can be composite in nature).
- We say here “ x determines y ” or “ y is functionally dependent on x ”
 $X \rightarrow Y$ does not imply $Y \rightarrow X$
- If the value of an attribute “Marks” is known then the value of an attribute “Grade” is determined since $\text{Marks} \rightarrow \text{Grade}$
- Types of functional dependencies:
 - Full Functional dependency
 - Partial Functional dependency
 - Transitive dependency

(Self Learning)

Consider the following Relation

REPORT (**STUDENT#, COURSE#**, **CourseName**, **IName**, **Room#**, **Marks**, **Grade**)

- **STUDENT#** - Student Number
 - **COURSE#** - Course Number
 - **CourseName** - Course Name
 - **IName** - Name of the Instructor who delivered the course
 - **Room#** - Room number which is assigned to respective Instructor
 - **Marks** - Scored in Course COURSE# by Student STUDENT#
 - **Grade** - obtained by Student STUDENT# in Course COURSE#
-

(Self Learning)

Functional Dependencies- From the previous example


- **STUDENT# COURSE# \rightarrow Marks**
 - **COURSE# \rightarrow CourseName,**
 - **COURSE# \rightarrow IName (Assuming one course is taught by one and only one Instructor)**
 - **IName \rightarrow Room# (Assuming each instructor has his/her own and non-shared room)**
 - **Marks \rightarrow Grade**
-

(Self Learning)

Dependency diagram

Report(S#, C#, SName, CTitle, LName, Room#, Marks, Grade)

- $S\# \rightarrow SName$
- $C\# \rightarrow CTitle,$
- $C\# \rightarrow LName$
- $LName \rightarrow Room\#$
- $C\# \rightarrow Room\#$
- $S\# C\# \rightarrow Marks$
- $Marks \rightarrow Grade$
- $S\# C\# \rightarrow Grade$



Assumptions:

- Each course has only one lecturer and each lecturer has a room.
- Grade is determined from Marks.


(Self Learning)

Full dependencies

X and Y are attributes.

X Functionally determines Y

Note: Subset of X should not functionally determine Y




(Self Learning)

Partial dependencies

X and Y are attributes.

Attribute Y is partially dependent on the attribute X only if it is dependent on a sub-set of attribute X.



(Self Learning)


Transitive dependencies

X Y and Z are three attributes.

$X \rightarrow Y$

$Y \rightarrow Z$

$\Rightarrow X \rightarrow Z$



Normal Forms - 1st Normal Form (1NF)

- A table (relation) is in 1NF if
 - There are no duplicated rows in the table.
 - Each cell is single-valued (i.e., there are no repeating groups or arrays).
 - Entries in a column (attribute, field) are of the same kind.
-

Recall

Student_Course_Result Table

Student_Details			Course_Details				Results		
101	Davis	11/4/1986	M4	Applied Mathematics	Basic Mathematics	7	11/11/2004	82	A
102	Daniel	11/6/1987	M4	Applied Mathematics	Basic Mathematics	7	11/11/2004	62	C
101	Davis	11/4/1986	116	American History		4	11/22/2004	79	D
103	Sandra	10/2/1988	C3	Bio Chemistry	Basic Chemistry	11	11/16/2004	65	B
104	Evelyn	2/22/1986	R3	Botany		8	11/26/2004	77	B
102	Daniel	11/6/1987	P3	Nuclear Physics	Basic Physics	13	11/12/2004	68	B
105	Susan	8/31/1985	P3	Nuclear Physics	Basic Physics	13	11/12/2004	89	A
103	Sandra	10/2/1988	R4	Zoology		5	11/27/2004	51	D
105	Susan	8/31/1985	116	American History		4	11/22/2004	87	A
104	Evelyn	2/22/1986	M4	Applied Mathematics	Basic Mathematics	7	11/11/2004	65	B

Normal Forms – 2nd Normal Form (2NF)

- *A Relation is said to be in Second Normal Form if and only if :*
 - *It is in the First normal form, and*
 - *No partial dependency exists between non-key attributes and key attributes.*

- An attribute of a relation R that belongs to any key of R is said to be a prime attribute and that which doesn't is a **non-prime attribute**

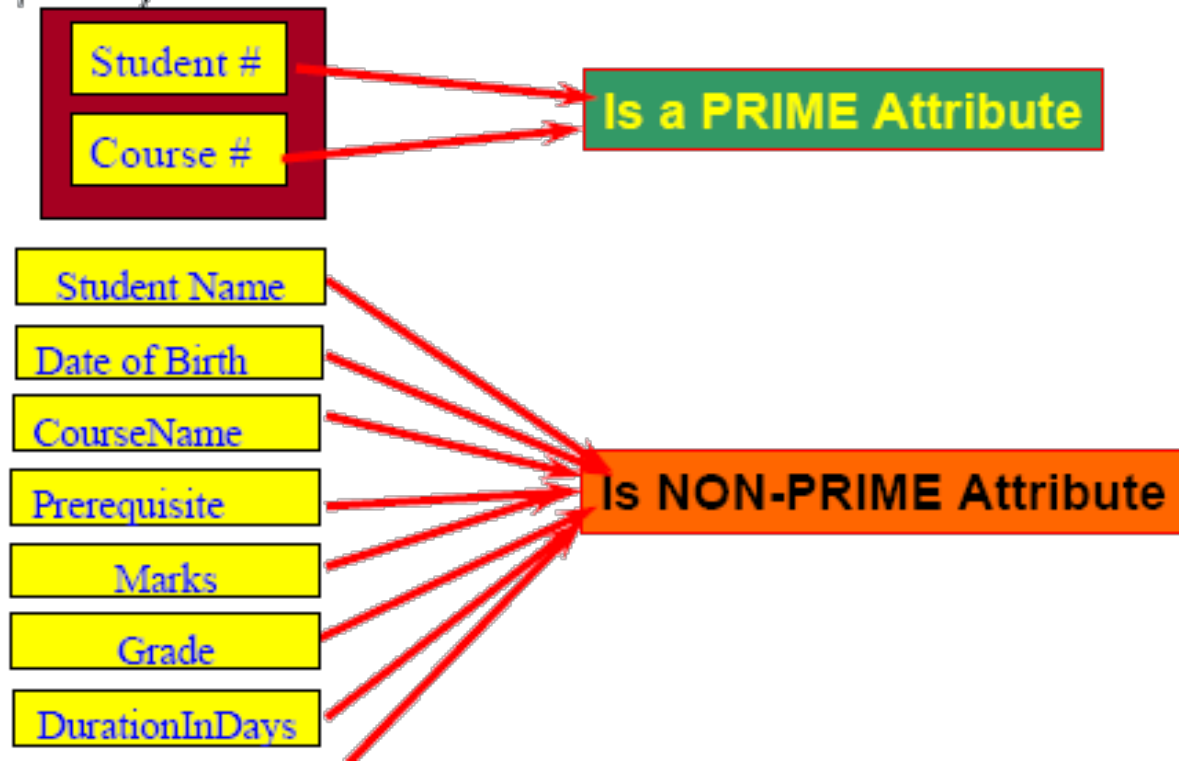
To make a table 2NF compliant, we have to remove all the partial dependencies

Note : - All partial dependencies are eliminated

Prime Vs Non-prime attributes

- An attribute of a relation R that belongs to any key of R is said to be a **prime attribute** and that which doesn't is a **non-prime attribute**

Report(S#,C#,StudentName,DateOfBirth,CourseName,PreRequisite,DurationInDays,DateOfExam,Marks,Grade)



Normal Forms – 2nd Normal Form (2NF)

- STUDENT# is key attribute for Student,
- COURSE# is key attribute for Course
- STUDENT# COURSE# together form the composite key attributes for Results relationship.
- Other attributes like StudentName (Student Name), DateofBirth, CourseName, PreRequisite, DurationInDays, DateofExam, Marks and Grade are non-key attributes.

To make this table 2NF compliant, we have to remove all the partial dependencies.


Student #, Course# -> Marks, Grade

Student# -> StudentName, DOB,

Course# -> CourseName, Prerequisite, DurationInDays

Course# -> Date of Exam

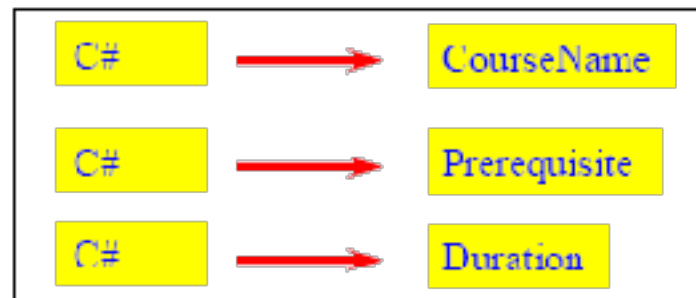
Normal Forms – 2nd Normal Form (2NF)




**Fully Functionally
dependent on composite
Candidate key**



Partial Dependency



Partial Dependency



Partial Dependency

Normal Forms – 2nd Normal Form (2NF)

STUDENT TABLE

Student#	StudentName	DateofBirth
101	Davis	04-Nov-1986
102	Daniel	06-Nov-1987
103	Sandra	02-Oct-1988
104	Evelyn	22-Feb-1986
105	Susan	31-Aug-1985
106	Mike	04-Feb-1987
107	Juliet	09-Nov-1986
108	Tom	07-Oct-1986

COURSE TABLE

Course#	Course Name	Pre Requisite	Duration InDays
M1	Basic Mathematics		11
M4	Applied Mathematics	M1	7
H6	American History		4
C1	Basic Chemistry		5
C3	Bio Chemistry	C1	11
B3	Botany		8
P1	Basic Physics		8

Normal Forms – 2nd Normal Form (2NF)

Student-Marks

Student#	Course#	Marks	Grade
101	M4	82	A
102	M4	62	C
101	H6	79	B
103	C3	65	B
104	B3	77	B
102	P3	68	B
105	P3	89	A
103	B4	54	D
105	H6	87	A
104	M4	65	B

Exam-Date


Course#	DateOfExam
M4	11-Nov-04
H6	22-Nov-04
C3	16-Nov-04
B3	26-Nov-04
P3	12-Nov-04
B4	27-Nov-04

Normal Forms – 3rd Normal Form (3NF)

A relation R is said to be in the Third Normal Form (3NF) if and only if

- It is in 2NF and*
- No transitive dependency exists between non-key attributes and key attributes.*

- STUDENT# and COURSE# are the key attributes.
- All other attributes, except grade are non-partially, non-transitively dependent on key attributes.
- **Student#, Course# -> Marks**
- **Marks -> Grade**



Note : - All transitive dependencies are eliminated

Normal Forms – 3rd Normal Form (3NF)

3NF Tables

Student#	Course#	Marks
101	M4	82
102	M4	62
101	H6	79
103	C3	65
104	B3	77
102	P3	68
105	P3	89
103	B4	54
105	H6	87
104	M4	65

Third Normal Form – Tables in 3 NF

MARKSGRADE TABLE		
UpperBound	LowerBound	Grade
100	95	A+
94	85	A
84	70	B
69	65	B-
64	55	C
54	45	D

MySQL - Session 1

Objectives

- After completing this lesson, you should be able to:
 - Describe the main database objects
 - Create tables
 - Describe the data types that can be used when specifying column definition
 - Alter table definitions
 - Drop, rename, and truncate tables
-

Database Objects in MySQL

Object	Description
Table	Basic unit of storage; composed of rows and columns
View	Logically represents subsets of data from one or more tables
Index	Improves the performance of some queries

Naming Rules For Identifiers

- Identifiers are names of table, view, index, column, constraint etc.
 - An identifier may be quoted or unquoted
 - Permitted characters in unquoted identifiers:
 - ASCII: [0-9, a-z, A-Z \$ _] (basic Latin letters, digits 0-9, dollar, underscore)
 - Extended: U+0080 .. U+FFFF
 - Identifiers may begin with a digit but unless quoted may not consist solely of digits.
 - Database, table, and column names cannot end with space characters.
 - Database and table names cannot contain “/”, “\”, “.”, or characters that are not permitted in file names.
 - The identifier quote character is the backtick (“`”)
-

Table creation

```
CREATE TABLE [IF NOT EXISTS] tbl_name(  
  create_definition,  
  ...  
);
```

create_definition may be one or more of the following:

- *col_name column_definition*
- [CONSTRAINT [*constraint_name*]] PRIMARY KEY (*col_name*)
- [CONSTRAINT [*constraint_name*]] UNIQUE (*col_name*)
- [CONSTRAINT [*constraint_name*]] FOREIGN KEY (*col_name*)
 reference_definition
- CHECK (*expr*)

Note: the CHECK constraint is parsed, but not implemented by the MySQL engines.

For a complete syntax/features refer:

<http://dev.mysql.com/doc/refman/5.1/en/create-table.html>

Table creation

To define a column in a table, use the following syntax:

col_name *column_definition*

where

column_definition:

data_type

[NOT NULL | NULL]

[DEFAULT default_value]

[AUTO_INCREMENT]

[UNIQUE | PRIMARY KEY]

[reference_definition]

reference_definition:

REFERENCES tbl_name (index_col_name, ...)

[ON DELETE *reference_option*]

[ON UPDATE *reference_option*]

reference_option:

RESTRICT | CASCADE | SET NULL | NO ACTION

Examples of column definition

- age int
- emp_id int primary key
- emp_id int primary key auto_increment
- city varchar(20) not null default 'Bangalore'
- email_id varchar(50) not null unique

- dept_id int references tbl_departments (dept_id)
- dept_id int references tbl_departments(dept_id) on delete cascade
- age int check (age>0)

Note: The last 3 statements will not give any syntax errors, but do not work as expected.

Datatypes

- Datatypes in MySQL are broadly divided into the following categories:
 - Numeric (integers and real)
 - Date and time
 - String

For more details about the datatypes in MySQL, please visit

<http://dev.mysql.com/doc/refman/5.1/en/data-types.html>

Datatypes – Numeric (integers)

Type	Storage (Bytes)	Minimum Value (Signed/Unsigned)	Maximum Value Signed/Unsigned)
TINYINT	1	-128	127
		0	255
SMALLINT	2	-32768	32767
		0	65535
MEDIUMINT	3	-8388608	8388607
		0	16777215
INT	4	-2147483648	2147483647
		0	4294967295
BIGINT	8	-9223372036854775808	9223372036854775807
		0	18446744073709551615

Datatypes – Numeric (real)

- Floating Point
 - FLOAT (uses 4 bytes to store the data)
 - DOUBLE (uses 8 bytes to store the data)
 - Fixed Point
 - DECIMAL (precision, scale)
-

Datatypes – Date and time

- DATE
 - Used for values with a date part but no time part.
 - MySQL retrieves and displays DATE values in 'YYYY-MM-DD' format.
 - The supported range is '1000-01-01' to '9999-12-31'.
 - DATETIME
 - Used for values that contain both date and time parts.
 - MySQL retrieves and displays DATETIME values in 'YYYY-MM-DD HH:MM:SS' format.
 - The supported range is '1000-01-01 00:00:00' to '9999-12-31 23:59:59'
 - TIMESTAMP
 - Used for values that contain both date and time parts.
 - The supported range is **'1970-01-01 00:00:01' to '2038-01-19 03:14:07'**
-

Datatypes – Date and time

- TIME
 - MySQL retrieves and displays TIME values in 'HH:MM:SS' format
 - May range from '-838:59:59' to '838:59:59'
 - The hours part may also represent elapsed time or a time interval between two events
-

Datatypes - String

- CHAR
 - VARCHAR
 - BLOB
 - TEXT
 - ENUM
-

Datatypes – String (CHAR)

- CHAR
 - Length is fixed
 - Length can range from 0 to 255 characters
 - When CHAR values are stored, they are right-padded with spaces to the specified length.
 - When CHAR values are retrieved, trailing spaces are removed unless the PAD_CHAR_TO_FULL_LENGTH SQL mode is enabled.
 - SET SQL_MODE='PAD_CHAR_TO_FULL_LENGTH'
-

Datatypes – String (VARCHAR)

- VARCHAR
 - Values in VARCHAR columns are variable-length strings.
 - The length can be specified as a value from 0 to 65535.
-

Datatypes – String (BLOB, TEXT)

- BLOB (TINYBLOB, BLOB, MEDIUMBLOB, LONGBLOB)
 - Binary Large Object that can hold a variable amount of data
 - Usually used for storing non-textual file content such as images, audio or video
 - BLOB values are treated as binary strings (byte strings).
 - They have no character set, and sorting and comparison are based on the numeric values of the bytes in column values.
 - Maximum size of a BLOB field is 64 KB.
 - If you need to store more than the above mentioned size, then use MEDIUMBLOB (16 MB) or LONGBLOB (4GB)
 - TEXT (TINYTEXT, TEXT, MEDIUMTEXT, LONGTEXT)
 - Same as BLOB, except that it can hold text data (non-binary string)
 - Sizes are similar to their BLOB counterparts
-

Datatypes – String (ENUM)

- An ENUM is a string object with a value chosen from a list of permitted values that are enumerated explicitly in the column specification at table creation time.

```
CREATE TABLE QUESTIONS (  
  ID INT PRIMARY KEY,  
  QUESTION TEXT,  
  DIFFICULTY_LEVEL ENUM ('VERY EASY', 'EASY',  
    'MODERATE', 'DIFFICULT', 'VERY DIFFICULT')  
);
```

- By default, if an invalid value is inserted for this column, MySQL will insert a NULL
 - This can be controlled by setting the SQL_MODE to 'TRADITIONAL', which results in an error, if an invalid value is inserted
-

Datatypes – String (ENUM)

```
mysql> SET SQL_MODE='';
Query OK, 0 rows affected (0.00 sec)

mysql> INSERT INTO QUESTIONS VALUES (1, 'WHAT IS THE CAPITAL CITY OF INDIA', 'SO EASY');
Query OK, 1 row affected, 1 warning (0.04 sec)

mysql> SELECT * FROM QUESTIONS;
+-----+-----+-----+
| ID | QUESTION | DIFFICULTY_LEVEL |
+-----+-----+-----+
| 1 | WHAT IS THE CAPITAL CITY OF INDIA | |
+-----+-----+-----+
1 row in set (0.00 sec)

mysql> SET SQL_MODE='TRADITIONAL';
Query OK, 0 rows affected (0.00 sec)

mysql> INSERT INTO QUESTIONS VALUES (2, 'WHAT IS THE CAPITAL CITY OF KARNATAKA', 'SO EASY');
ERROR 1265 (01000): Data truncated for column 'DIFFICULTY_LEVEL' at row 1
mysql> SELECT * FROM QUESTIONS;
+-----+-----+-----+
| ID | QUESTION | DIFFICULTY_LEVEL |
+-----+-----+-----+
| 1 | WHAT IS THE CAPITAL CITY OF INDIA | |
+-----+-----+-----+
1 row in set (0.00 sec)

mysql>
```

With SQL_MODE=""
The record was inserted with
NULL value for
DIFFICULTY_LEVEL

With SQL_MODE='TRADITIONAL'
The record was not inserted.

Table - examples

```
CREATE TABLE BOOKS(  
    BOOK_ID INT AUTO_INCREMENT PRIMARY KEY,  
    TITLE VARCHAR(30) NOT NULL,  
    PRICE DOUBLE  
);
```

Table - examples

```
CREATE TABLE DEPTS (  
    DEPT_ID INT,  
    DNAME VARCHAR(20) NOT NULL,  
    PRIMARY KEY (DEPT_ID)  
);
```

```
CREATE TABLE EMPS (  
    EMP_ID INT PRIMARY KEY,  
    ENAME VARCHAR(20),  
    DEPT_NO INT,  
    SALARY DECIMAL(10, 2),  
    FOREIGN KEY (DEPT_NO) REFERENCES DEPTS(DEPT_ID)  
        ON DELETE SET NULL  
);
```

Constraints

- The MySQL Server uses constraints to prevent invalid data entry into tables.
 - You can use constraints to do the following:
 - Enforce rules on the data in a table whenever a row is inserted, updated, or deleted from that table.
 - Prevent the deletion of a table if there are dependencies from other tables
 - Constraints to a table can be added during the table creation itself, or later using the ALTER TABLE command
 - Constraints can be given at table or column level
-

Constraints

- NOT NULL
 - Does not allow blanks for a given column
 - Note that ' or "" is not same as NULL
 - UNIQUE
 - Does not allow duplicates
 - MySQL allows more than one NULL entry to a column with UNIQUE constraint (unlike Oracle)
 - Since NULL is not a value, repetition of NULL values for a given column is not a violation of this constraint. To avoid this, use UNIQUE with NOT NULL constraint.
-

Constraints

- PRIMARY KEY
 - One or more columns together identify a row uniquely.
 - Only one PRIMARY KEY is allowed per table
 - If more than one column are involved, then the constraint must be specified as a table level constraint.
 - This kind of primary key is also known as COMPOSITE KEY
 - FOREIGN KEY
 - Establishes a PARENT-CHILD relationship between two tables
 - The table being referenced is considered as a PARENT table
 - The table which is referencing (where the constraint is specified), is considered as a CHILD table
-

Constraints - examples

```
create table person (  
  id int,  
  name varchar(20) not null,  
  city varchar(20),  
  constraint PK_PERSON primary key (id));
```

Valid inserts:

```
insert into person values(1, 'SCOTT', 'DALLAS');  
insert into person (id, name) values (2, 'MILLER');  
insert into person values (3, 'ALLEN', NULL);
```

Invalid inserts:

```
insert into person values (2, 'SMITH', 'CHICAGO');  
    → duplicate value for primary key  
insert into person (id, city) values (4, 'BANGALORE');  
    → name can not contain null  
insert into person values (4, NULL, 'MUMBAI');  
    → name can not contain null
```

Constraints - examples

Column level constraint:

```
CREATE TABLE BOOKS(  
    BOOK_ID INT PRIMARY KEY,  
    TITLE VARCHAR(30) NOT NULL,  
    PRICE DOUBLE  
);
```

While specifying the column-level constraint, a constraint name can not be given.

Table level constraint:

```
CREATE TABLE BOOKS(  
    BOOK_ID INT,  
    TITLE VARCHAR(30) NOT NULL,  
    PRICE DOUBLE,  
    CONSTRAINT PK_BOOKS PRIMARY KEY (BOOK_ID)  
);
```

Hence, it is a good practice to use the table level constraint specification.

Constraints - examples

Column level constraint:

```
CREATE TABLE EMPS (  
    EMP_ID INT PRIMARY KEY,  
    ENAME VARCHAR(20),  
    DEPT_NO INT REFERENCES DEPTS(DEPT_ID),  
    SALARY DECIMAL(10, 2)  
);
```

Table level constraint:

```
CREATE TABLE EMPS (  
    EMP_ID INT PRIMARY KEY,  
    ENAME VARCHAR(20),  
    DEPT_NO INT,  
    SALARY DECIMAL(10, 2),  
    CONSTRAINT FK_EMPS_DEPT  
    FOREIGN KEY (DEPT_NO) REFERENCES DEPTS(DEPT_ID)  
);
```

The **column-level foreign key** constraint specification is successfully parsed by MySQL, but may not enforce the constraint.

Hence, always use the table level constraint specification for **foreign keys**.

Constraints – Example on foreign keys

```
CREATE TABLE CATEGORIES (  
  ID INT PRIMARY KEY,  
  CATEGORY_NAME VARCHAR(20));
```

```
INSERT INTO CATEGORY VALUES(10, 'COMPUTERS');  
INSERT INTO CATEGORY VALUES(20, 'STATIONARIES');
```

```
CREATE TABLE PRODUCTS (  
  ID INT PRIMARY KEY,  
  PRODUCT_NAME VARCHAR(20),  
  CATEGORY_ID INT,  
  CONSTRAINT FK_PRD_CAT  
  FOREIGN KEY (CATEGORY_ID) REFERENCES CATEGORIES(ID)  
);
```

```
INSERT INTO PRODUCTS VALUES (1, 'KEYBOARD', 10); -- Works fine  
INSERT INTO PRODUCTS VALUES (2, 'PURSE', 50); -- Fails, No Parent Key
```

```
DELETE FROM CATEGORIES WHERE ID = 20; -- Works fine  
DELETE FROM CATEGORIES WHERE ID = 10; -- Fails, child records exist
```

Constraints – Foreign key UPDATE/DELETE options

- In a parent-child relationship using the foreign key constraint, when a record in the parent table is deleted,
 - it will be successful if there is no child record corresponding to the parent record being deleted.
 - it will fail if there is a child record corresponding to the parent record being deleted.
 - This can be controlled using the ON DELETE option of the constraint.

```
ALTER TABLE PRODUCTS DROP FOREIGN KEY FK_PRD_CAT;
```

```
ALTER TABLE PRODUCTS  
    ADD CONSTRAINT FOREIGN KEY FK_PRD_CAT (CATEGORY_ID)  
    REFERENCES CATEGORIES (ID)  
    ON DELETE CASCADE;
```

Constraints – Foreign key DELETE/UPDATE options

- The different values for ON DELETE option are:
 - CASCADE
 - Dangerous, since this will delete all the corresponding child records.
 - SET NULL
 - Safe, allows the parent record to be deleted, while preserving the child records.
 - For this to work, the foreign key column should not have NOT NULL constraint on it.
 - RESTRICT
 - Default and should be preferred in most cases to ensure that the data is not accidentally deleted.
 - The same values apply for 'ON UPDATE' option of the constraint.
-

Altering the table structure

- Alter table allows us to do the following:
 - Change the name of the column
 - Change the column definition
 - Add a new column
 - Delete an existing column
 - Add/delete primary key
 - Add/delete foreign key
 - Rename a table
-

Syntax: Alter table

- ALTER TABLE tbl_name
 - [alter_specification [, alter_specification] ...]
 - Where alter_specification can be:
 - ADD [COLUMN] col_name column_definition
 - ADD [CONSTRAINT [con_name]] PRIMARY KEY (col_name)
 - ADD [CONSTRAINT [con_name]] UNIQUE (col_name)
 - ADD [CONSTRAINT [con_name]] FOREIGN KEY (col_name) reference_definition
 - ALTER [COLUMN] col_name {SET DEFAULT literal | DROP DEFAULT}
 - CHANGE [COLUMN] old_col_name new_col_name column_definition
 - MODIFY [COLUMN] col_name column_definition
 - ..contd
-

Syntax: Alter table

- ..cont.
 - DROP [COLUMN] col_name
 - DROP PRIMARY KEY
 - DROP FOREIGN KEY fk_symbol
 - RENAME [TO] new_tbl_name
-

Alter table - examples

```
ALTER TABLE QUESTIONS  
ADD COLUMN MARKS INT;
```

```
ALTER TABLE QUESTIONS  
CHANGE COLUMN MARKS POINTS INT;
```

```
ALTER TABLE QUESTIONS  
DROP PRIMARY KEY;
```

```
ALTER TABLE QUESTIONS  
ADD PRIMARY KEY(ID);
```

```
ALTER TABLE EMPLOYEES  
ALTER CITY SET DEFAULT 'BANGALORE'
```

```
ALTER TABLE EMPLOYEES  
ALTER CITY DROP DEFAULT;
```

Data Manipulation Language

- A DML statement corresponds to
 - adding a new row to a table
 - modifying existing rows in a table
 - deleting existing rows from a table
 - A transaction consists of a collection of DML statements that form a logical unit of work.
-

The INSERT statement - Syntax

```
INSERT [INTO] tbl_name [(col_name,...)]  
    {VALUES | VALUE} ({expr | DEFAULT},...),(...),...
```

The above syntax allows us to insert one or more rows to a given table.

Examples:

```
INSERT CATEGORIES VALUES (10, 'COMPUTERS');
```

```
INSERT INTO CATEGORIES VALUE (20, 'STATIONARIES');
```

```
-- Insert multiple rows
```

```
INSERT CATEGORIES VALUES (30, 'BEVERAGES'), (40, 'CONDIMENTS');
```

```
-- Selected columns only
```

```
INSERT PRODUCTS (ID, PRODUCT_NAME)
```

```
VALUES (1, 'KEYBOARD'), (2, 'OPTICAL MOUSE'), (3, 'SCANNER');
```

In the above example, the omitted columns will be inserted with either NULL or DEFAULT (if specified in column definition)

The INSERT statement

- While inserting text or date/time values, you can use either single quote or double quote.
 - Note: Other database softwares may allow only single quotes as text qualifiers and hence, it is a good practice to use the same in MySQL also.
 - You can use functions during a record insertion.
 - For example, to insert a registration date as current date, you may use
`INSERT INTO MEMBERS VALUES (9812, 'SCOTT', SYSDATE(), 'DALLAS');`
 - While inserting date values, use the YYYY-MM-DD format
 - `INSERT INTO MEMBERS VALUES (1988, 'MIKE', '2012-07-22', 'UTAH');`
-

The INSERT statement

- You can also copy values from another table using the INSERT statement:
 - ```
INSERT INTO TOP_10_EMPS
 SELECT EMPNO, ENAME, SALARY FROM EMPLOYEES
 ORDER BY SALARY DESC LIMIT 10
```
  - The column types in the query should match with the column types in the target table.
  - Do not use the VALUE/VALUES during this operation.
-

# Modifying existing rows – UPDATE statement

- You can change the value of one or more columns of one or more rows in a table using the UPDATE statement.

```
UPDATE table_name
SET col_name1={expr1|DEFAULT} [, col_name2={expr2|DEFAULT}]
...
[WHERE where_condition]
[ORDER BY column_name [{ASC|DESC}]]
[LIMIT row_count]
```

- If you do not specify the WHERE clause, then all the rows in the table will be updated
-



# UPDATE - Examples

Increase the product price by 10% for all products:

```
UPDATE PRODUCTS
```

```
 SET PRICE = PRICE + PRICE * 10 / 100;
```

```
UPDATE PRODUCTS
```

```
 SET PRICE = PRICE * 1.1;
```

Increase the price by \$5 for those products which are priced below \$10:

```
UPDATE PRODUCTS
```

```
 SET PRICE = PRICE + 5
```

```
 WHERE PRICE < 10;
```

Convert the CITY, STATE, and COUNTRY into upper case letters:

```
UPDATE CUSTOMERS
```

```
 SET CITY=UPPER('CITY'),
```

```
 STATE = UPPER('STATE'),
```

```
 COUNTRY = UPPER('COUNTRY');
```

---

# UPDATE - Examples

Change the CITY of the CUSTOMER whose ID is 1245 to the DEFAULT value:

```
UPDATE CUSTOMERS
 SET CITY = DEFAULT
 WHERE CUSTOMER_ID = 1245;
```

Change the DESIGNATION of TOP 3 salaried employees as 'DIRECTOR':

```
UPDATE EMPLOYEES
 SET DESIGNATION = 'DIRECTOR'
 ORDER BY SALARY DESC
 LIMIT 3;
```

---

# UPDATE - Examples

Change employee 114's job and salary to match that of employee 205.

```
UPDATE EMPLOYEES AS E1
 JOIN EMPLOYEES AS E2
 ON E2.EMPLOYEE_ID=205
 SET E1.JOB_ID = E2.JOB_ID,
 E1.SALARY = E2.SALARY
 WHERE
 E1.EMPLOYEE_ID = 114;
```

The source data and the corresponding condition.

The columns whose values need to be updated.

The rows in the target table to be updated.

In other database softwares, such as Oracle, the same thing can be achieved using a subquery as shown below:

```
UPDATE EMPLOYEES
SET JOB = (SELECT JOB FROM EMPLOYEES WHERE EMPLOYEE_ID = 205),
SALARY = (SELECT SALARY FROM EMPLOYEES WHERE EMPLOYEE_ID = 205)
WHERE EMPLOYEE_ID = 114;
```

However, in MySQL, this is not allowed.

**You can not modify the same table which you are using in the SELECT part.**

---

# Updating rows – Integrity constraint violations

Trying to change the parent row's primary key value, while there are child records existing for this parent row:

```
UPDATE CATEGORIES SET ID=15 WHERE ID=10;
```

ERROR 1451 (23000): Cannot delete or update a parent row: a foreign key constraint fails (`mydb`.`products`, CONSTRAINT `fk\_prd\_cat` FOREIGN KEY (`CATEGORY\_ID`) REFERENCES `categories` (`ID`) ON DELETE CASCADE)

Trying to change the child row's foreign key value, while the new foreign key value is not present in the parent table:

```
UPDATE PRODUCTS SET CATEGORY_ID=15 WHERE ID = 3;
```

ERROR 1452 (23000): Cannot add or update a child row: a foreign key constraint fails (`mydb`.`products`, CONSTRAINT `fk\_prd\_cat` FOREIGN KEY (`CATEGORY\_ID`) REFERENCES `categories` (`ID`) ON DELETE CASCADE)

---

# Deleting Rows from a Table - the DELETE Statement

- You can remove existing rows from a table by using the DELETE statement.

```
DELETE [FROM] <table_name>
[WHERE <condition>]
```

## Deleting Rows from a Table - the DELETE Statement

```
mysql> select category_id, category_name from categories;
```

| category_id | category_name  |
|-------------|----------------|
| 1           | Beverages      |
| 2           | Condiments     |
| 3           | Confections    |
| 4           | Dairy Products |
| 5           | Grains/Cereals |
| 6           | Meat/Poultry   |
| 7           | Produce        |
| 8           | Seafood        |
| 15          | Softdrinks     |
| 17          | Electronics    |

```
10 rows in set (0.00 sec)
```

```
mysql> delete from categories where category_id in (15, 17);
Query OK, 2 rows affected (0.09 sec)
```

Specific rows are deleted if you specify the **WHERE** clause.

```
mysql> select category_id, category_name from categories;
```

| category_id | category_name  |
|-------------|----------------|
| 1           | Beverages      |
| 2           | Condiments     |
| 3           | Confections    |
| 4           | Dairy Products |
| 5           | Grains/Cereals |
| 6           | Meat/Poultry   |
| 7           | Produce        |
| 8           | Seafood        |

```
8 rows in set (0.00 sec)
```

# Deleting Rows from a Table - the DELETE Statement

All rows in the table are deleted  
if you omit the WHERE clause.

```
mysql> delete from categories;
Query OK, 8 rows affected (0.00 sec)

mysql> select category_id, category_name from categories;
Empty set (0.00 sec)
```

---

# Database Transactions

- A transaction comprises a unit of work performed within a database management system against a database, and treated in a coherent and reliable way independent of other transactions.
  - Transactions in a database environment have two main purposes:
  - A database transaction, by definition, must be atomic, consistent, isolated and durable.
    - Database practitioners often refer to these properties of database transactions using the acronym ACID.
-



# Database Transactions

- Transaction begins when the first DML SQL statement is executed.
  - Ends with one of the following events:
    - A COMMIT or ROLLBACK statement is issued
    - A DDL statement executes (automatic commit)
    - The system crashes
-

# Advantages of COMMIT and ROLLBACK Statements

- With COMMIT and ROLLBACK statements, you can:
    - Ensure data consistency
    - Preview data changes before making changes permanent
    - Group logically related operations
-

# Controlling Transactions



# Implicit Transaction Processing

- An automatic commit occurs under the following circumstances:
    - DDL statement is issued
    - The autocommit option is turned on
      - `SET AUTOCOMMIT=1;`
  - By default, the autocommit is turned on by MySQL, and if you want it to be turned off, then issue this command:
    - `SET AUTOCOMMIT=0;`
-

# Rolling back the data changes

- Discard all pending changes by using the ROLLBACK statement:
    - Data changes are undone.
    - Previous state of the data is restored.
    - Locks on the affected rows are released.
-

# Rolling back the data changes

```
mysql> set autocommit=0;
Query OK, 0 rows affected (0.00 sec)

mysql> insert into categories values (60, 'SEA FOODS');
Query OK, 1 row affected (0.00 sec)

mysql> select * from categories;
+----+-----+
| ID | CATEGORY_NAME |
+----+-----+
10	COMPUTERS
20	STATIONARIES
30	BEVERAGES
40	CONDIMENTS
50	CONFECTIONS
60	SEA FOODS
+----+-----+
6 rows in set (0.00 sec)

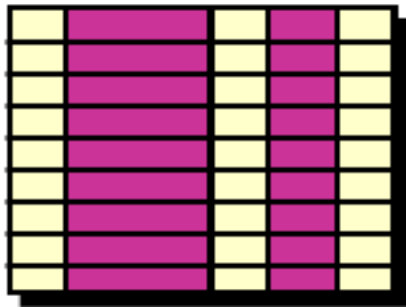
mysql> rollback;
Query OK, 0 rows affected (0.04 sec)

mysql> select * from categories;
+----+-----+
| ID | CATEGORY_NAME |
+----+-----+
10	COMPUTERS
20	STATIONARIES
30	BEVERAGES
40	CONDIMENTS
50	CONFECTIONS
+----+-----+
5 rows in set (0.00 sec)
```

# Queries

# Different operations on a table

**Projection**

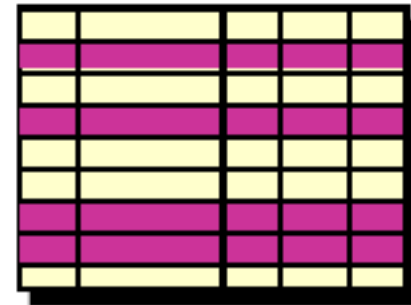


A diagram illustrating the Projection operation. It shows a 10x5 grid representing Table 1. The second and fourth columns are highlighted in pink, indicating they are the selected attributes for the projection.

|  |  |  |  |  |
|--|--|--|--|--|
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |

**Table 1**

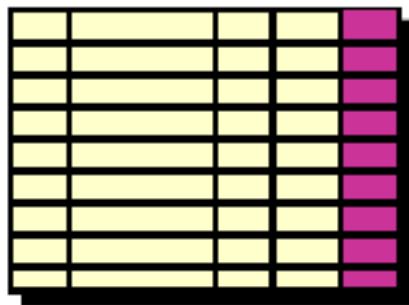
**Selection**



A diagram illustrating the Selection operation. It shows a 10x5 grid representing Table 1. The first, third, and fifth columns are highlighted in pink, indicating they are the selected rows.

|  |  |  |  |  |
|--|--|--|--|--|
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |

**Table 1**

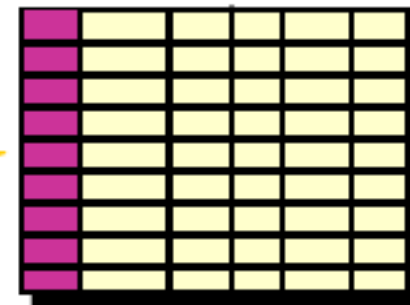


A diagram showing the result of the Projection operation. It is a 10x5 grid where only the second and fourth columns are present and highlighted in pink, representing the projected attributes.

|  |  |  |  |  |
|--|--|--|--|--|
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |

**Table 1**

**Join**



A diagram showing the result of the Join operation. It is a 10x5 grid where the first column is highlighted in pink, representing the joined attribute.

|  |  |  |  |  |
|--|--|--|--|--|
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |
|  |  |  |  |  |

**Table 2**



# SELECT - Syntax

```
SELECT [DISTINCT] col_name [, col_name...]
[FROM table_references]
[WHERE where_condition]
[GROUP BY {col_name} [HAVING where_condition]]
[ORDER BY {col_name} [ASC | DESC], ...]
[LIMIT row_count]
[OFFSET offset]
```

- SELECT identifies which columns to retrieve
- FROM identifies which tables to retrieve the data from
- WHERE identifies which rows to filter (or retrieve)
- GROUP BY identifies how to summarize the data
- ORDER BY identifies how the data need to be arranged
- LIMIT identifies how many rows to be retrieved within the result
- OFFSET identifies which record to start retrieving the result from

You can refer the following link for the detailed syntax and usage:

<http://dev.mysql.com/doc/refman/5.1/en/select.html>

---

# Selecting all columns

```
SELECT *
FROM SHIPPERS;
```

```
mysql> SELECT * FROM SHIPPERS;
+-----+-----+-----+
| SHIPPER_ID | COMPANY_NAME | PHONE |
+-----+-----+-----+
1	Speedy Express	(503) 555-9831
2	United Package	(503) 555-3199
3	Federal Shipping	(503) 555-9931
+-----+-----+-----+
3 rows in set (0.07 sec)
```

# Selecting specific columns (Projection)

```
SELECT CATEGORY_NAME, DESCRIPTION
FROM CATEGORIES;
```

```
+-----+-----+
| CATEGORY_NAME | DESCRIPTION |
+-----+-----+
Beverages	Soft drinks, coffees, teas, beers, and ales
Condiments	Sweet and savory sauces, relishes, spreads, and seasonings
Confections	Desserts, candies, and sweet breads
Dairy Products	Cheeses
Grains/Cereals	Breads, crackers, pasta, and cereal
Meat/Poultry	Prepared meats
Produce	Dried fruit and bean curd
Seafood	Seaweed and fish
+-----+-----+
8 rows in set (0.07 sec)
```

# Writing SQL Statements

- SQL statements are not case sensitive
  - SQL statements can be on one or more lines.
  - Keywords cannot be abbreviated or split across lines.
  - Clauses are usually placed on separate lines.
  - Indents are used to enhance readability.
-

# Arithmetic Expressions

- Create expressions with number and date data by using arithmetic operators.

| Operator | Description |
|----------|-------------|
| +        | Add         |
| -        | Subtract    |
| *        | Multiply    |
| /        | Divide      |
| %        | Modulus     |

---

# Using arithmetic operators

```
SELECT PRODUCT_NAME, UNIT_PRICE,
UNIT_PRICE * 0.5 AS PROFIT
FROM PRODUCTS LIMIT 10;
```

| PRODUCT_NAME                    | UNIT_PRICE | PROFIT   |
|---------------------------------|------------|----------|
| Chai                            | 18.0000    | 9.00000  |
| Chang                           | 19.0000    | 9.50000  |
| Aniseed Syrup                   | 10.0000    | 5.00000  |
| Chef Anton's Cajun Seasoning    | 22.0000    | 11.00000 |
| Chef Anton's Gumbo Mix          | 21.3500    | 10.67500 |
| Grandma's Boysenberry Spread    | 25.0000    | 12.50000 |
| Uncle Bob's Organic Dried Pears | 30.0000    | 15.00000 |
| Northwoods Cranberry Sauce      | 40.0000    | 20.00000 |
| Mishi Kobe Niku                 | 97.0000    | 48.50000 |
| Ikura                           | 31.0000    | 15.50000 |

10 rows in set (0.14 sec)

# Operator precedence

- Multiplication and division take priority over addition and subtraction.
  - Operators of the same priority are evaluated from left to right.
  - Parentheses are used to force prioritized evaluation and to clarify statements.
-

## Operator precedence - Example

```
SELECT PRODUCT_NAME,
UNIT_PRICE * UNITS_IN_STOCK - REORDER_LEVEL AS EXPR1,
UNIT_PRICE * (UNITS_IN_STOCK - REORDER_LEVEL) AS EXPR2
FROM PRODUCTS LIMIT 10;
```

| PRODUCT_NAME                    | EXPR1     | EXPR2     |
|---------------------------------|-----------|-----------|
| Chai                            | 692.0000  | 522.0000  |
| Chang                           | 298.0000  | -152.0000 |
| Aniseed Syrup                   | 105.0000  | -120.0000 |
| Chef Anton's Cajun Seasoning    | 1166.0000 | 1166.0000 |
| Chef Anton's Gumbo Mix          | 0.0000    | 0.0000    |
| Grandma's Boysenberry Spread    | 2975.0000 | 2375.0000 |
| Uncle Bob's Organic Dried Pears | 440.0000  | 150.0000  |
| Northwoods Cranberry Sauce      | 240.0000  | 240.0000  |
| Mishi Kobe Niku                 | 2813.0000 | 2813.0000 |
| Ikura                           | 961.0000  | 961.0000  |

10 rows in set (0.03 sec)



# NULL values

- A null is a value that is unavailable, unassigned, unknown, or inapplicable.
- A null is not the same as zero or a blank space.

```
SELECT FIRST_NAME, CITY,
REGION FROM EMPLOYEES;
```

| FIRST_NAME | CITY     | REGION |
|------------|----------|--------|
| Nancy      | Seattle  | WA     |
| Andrew     | Tacoma   | WA     |
| Janet      | Kirkland | WA     |
| Margaret   | Redmond  | WA     |
| Steven     | London   | NULL   |
| Michael    | London   | NULL   |
| Robert     | London   | NULL   |
| Laura      | Seattle  | WA     |
| Anne       | London   | NULL   |

# NULL values in expressions

- Arithmetic expressions containing a null value evaluate to null.

```
SELECT FIRST_NAME,
 CONCAT(CITY, ' ', REGION)
FROM EMPLOYEES;
```

| FIRST_NAME | CONCAT(CITY, ' ', REGION) |
|------------|---------------------------|
| Nancy      | Seattle WA                |
| Andrew     | Tacoma WA                 |
| Janet      | Kirkland WA               |
| Margaret   | Redmond WA                |
| Steven     | NULL                      |
| Michael    | NULL                      |
| Robert     | NULL                      |
| Laura      | Seattle WA                |
| Anne       | NULL                      |

# NULL values in expressions

- If any column value in an arithmetic expression is null, the result is null.
    - For example, if you attempt to perform division with zero, you get an error. However, if you divide a number by null, the result is a null or unknown.
-

# Defining a column alias

- A column alias:
    - Renames a column heading
    - Is useful with calculations
    - Immediately follows the column name
      - There can also be the optional '**AS**' keyword between the column name and alias
    - If the alias contains a space, then it must be quoted.
-

## Column alias - Example

```
SELECT CONCAT(FIRST_NAME, ' ', LAST_NAME) AS 'EMPLOYEE NAME', CITY
FROM EMPLOYEES;
```

| EMPLOYEE NAME    | CITY     |
|------------------|----------|
| Nancy Davolio    | Seattle  |
| Andrew Fuller    | Tacoma   |
| Janet Leverling  | Kirkland |
| Margaret Peacock | Redmond  |
| Steven Buchanan  | London   |
| Michael Suyama   | London   |
| Robert King      | London   |
| Laura Callahan   | Seattle  |
| Anne Dodsworth   | London   |

# Concatenation

- Unlike other DB softwares, MySQL does not use the double-pipe for concatenation.
  - For example, the following SQL statement would result in undesired output:  
`SELECT FIRST_NAME||' '||LAST_NAME AS 'EMPLOYEE NAME', CITY  
FROM EMPLOYEES;`
  - The reason for this is, MySQL considers double-pipe as 'logical OR' operator and not as concatenation operator.
  - You may use this operator by setting the SQL\_MODE to 'ANSI'

| EMPLOYEE NAME |  | CITY     |
|---------------|--|----------|
| Ø             |  | Seattle  |
| Ø             |  | Tacoma   |
| Ø             |  | Kirkland |
| Ø             |  | Redmond  |
| Ø             |  | London   |
| Ø             |  | London   |
| Ø             |  | London   |
| Ø             |  | Seattle  |
| Ø             |  | London   |

# Concatenation

```
mysql> SET SQL_MODE='ANSI';
Query OK, 0 rows affected (0.00 sec)

mysql> SELECT FIRST_NAME||' '||LAST_NAME AS 'EMPLOYEE NAME',
 -> CITY FROM EMPLOYEES;
+-----+-----+
| EMPLOYEE NAME | CITY |
+-----+-----+
Nancy Davolio	Seattle
Andrew Fuller	Tacoma
Janet Leverling	Kirkland
Margaret Peacock	Redmond
Steven Buchanan	London
Michael Suyama	London
Robert King	London
Laura Callahan	Seattle
Anne Dodsworth	London
+-----+-----+
9 rows in set (0.00 sec)
```

A better and preferred way of concatenating is to use the CONCAT function. You can supply multiple arguments to this function.

# Literal Character Strings

- A literal is a character, a number, or a date included in the SELECT list.
- Date and character literal values must be enclosed within single quotation marks.
- Each character string is output once for each row returned.

```
mysql> SELECT CONCAT(FIRST_NAME, ' is from ', CITY) AS EXPR
 -> FROM EMPLOYEES;
```

| EXPR                     |
|--------------------------|
| Nancy is from Seattle    |
| Andrew is from Tacoma    |
| Janet is from Kirkland   |
| Margaret is from Redmond |
| Steven is from London    |
| Michael is from London   |
| Robert is from London    |
| Laura is from Seattle    |
| Anne is from London      |

```
9 rows in set (0.00 sec)
```



# Restricting and Sorting Data

- The combination of SELECT and FROM retrieves all the rows.
- You can limit the number of rows based on a condition using the WHERE clause.

Retrieve all employees from 'London' city

```
mysql> SELECT FIRST_NAME, LAST_NAME FROM EMPLOYEES
-> WHERE CITY = 'LONDON';
```

| FIRST_NAME | LAST_NAME |
|------------|-----------|
| Steven     | Buchanan  |
| Michael    | Suyama    |
| Robert     | King      |
| Anne       | Dodsworth |

```
4 rows in set (0.00 sec)
```

# Character Strings and Dates

- Character strings and date values are enclosed in single quotation marks.
  - However, in the WHERE clause, strings and dates can be given in single or double quotes
- Character values are NOT case sensitive
  - Unless you are using MySQL on Unix/Linux platform
- The default date format is YYYY-MM-DD and datetime format is YYYY-MM-DD hh:mm:ss

```
mysql> SELECT LAST_NAME, FIRST_NAME, BIRTH_DATE, HIRE_DATE
-> FROM EMPLOYEES
-> WHERE LAST_NAME = 'KING';
```

| LAST_NAME | FIRST_NAME | BIRTH_DATE          | HIRE_DATE           |
|-----------|------------|---------------------|---------------------|
| King      | Robert     | 1960-05-29 00:00:00 | 1994-01-02 00:00:00 |

```
1 row in set (0.00 sec)
```

# Operators for WHERE clause

| Operator             | Description                                    |
|----------------------|------------------------------------------------|
| =                    | Equals to                                      |
| >                    | Greater than                                   |
| >=                   | Greater than or equals to                      |
| <                    | Less than                                      |
| <=                   | Less than or equals to                         |
| !=                   | Not equals to                                  |
| <>                   | Not equals to                                  |
| BETWEEN, NOT BETWEEN | Compare to a range of values                   |
| IN, NOT IN           | Match (or do not) any of the value in the list |
| LIKE, NOT LIKE       | Match (or do not) a character pattern          |
| IS NULL, IS NOT NULL | Compare with NULL value                        |

# WHERE - Examples

Retrieve the id, name, and price of all products that belong to category 1

```
SELECT PRODUCT_ID, PRODUCT_NAME, UNIT_PRICE
FROM PRODUCTS
WHERE CATEGORY_ID = 1;
```

Retrieve the id, name, price in stock for all products that are not in stock

```
SELECT PRODUCT_ID, PRODUCT_NAME, UNIT_PRICE
FROM PRODUCTS
WHERE UNITS_IN_STOCK = 0;
```

---

# WHERE - Examples

Retrieve the details of all the products cost less than \$20

```
SELECT * FROM PRODUCTS
WHERE UNIT_PRICE < 20;
```

Retrieve the details of all the products that are priced between \$20 and \$50

```
SELECT * FROM PRODUCTS
WHERE UNIT_PRICE BETWEEN 20 AND 50;
```

---

# WHERE - Examples

Retrieve the name and city of all customers who live in 'London', 'Paris', or 'Madrid'

```
SELECT CUSTOMER_NAME, CITY
FROM CUSTOMERS
WHERE CITY = 'LONDON' OR CITY = 'PARIS' OR CITY = 'MADRID';
```

```
SELECT CUSTOMER_NAME, CITY
FROM CUSTOMERS
WHERE CITY IN ('LONDON', 'PARIS', 'MADRID');
```

---

## WHERE - Examples

Retrieve the names of the suppliers who do not have a website.

```
SELECT COMPANY_NAME FROM SUPPLIERS
WHERE HOME_PAGE IS NULL;
```

Retrieve the names of the suppliers who have a website address along with their homepage.

```
SELECT COMPANY_NAME, HOME_PAGE FROM SUPPLIERS
WHERE HOME_PAGE IS NOT NULL;
```

---

# WHERE - Examples

Retrieve the name and city of the all customers whose name ends with 'EN'

```
SELECT CUSTOMER_NAME, CITY FROM CUSTOMERS
WHERE CUSTOMER_NAME LIKE '%EN';
```

Retrieve the name and joining date of all employees who joined in the year 1993

```
SELECT FIRST_NAME, LAST_NAME, HIRE_DATE FROM EMPLOYEES
WHERE HIRE_DATE LIKE '1993-%';
```

*A better and preferred method is using the year() function as shown below:*

```
SELECT FIRST_NAME, LAST_NAME, HIRE_DATE FROM EMPLOYEES
WHERE YEAR(HIRE_DATE) = 1993;
```

---



# Logical Operators

| Operator | Description                                      |
|----------|--------------------------------------------------|
| AND      | Returns TRUE if both conditions are true         |
| OR       | Returns TRUE if either of the conditions is true |
| NOT      | Returns TRUE if the following condition is false |

---

# Using the AND operator

- AND requires both conditions to be true

```
SELECT PRODUCT_NAME, CATEGORY_ID, UNIT_PRICE
FROM PRODUCTS
WHERE UNIT_PRICE BETWEEN 10 AND 50
AND
CATEGORY_ID = 1;
```

| PRODUCT_NAME              | CATEGORY_ID | UNIT_PRICE |
|---------------------------|-------------|------------|
| Chai                      | 1           | 18.0000    |
| Chang                     | 1           | 19.0000    |
| Sasquatch Ale             | 1           | 14.0000    |
| Steeleye Stout            | 1           | 18.0000    |
| Chartreuse verte          | 1           | 18.0000    |
| Ipoh Coffee               | 1           | 46.0000    |
| Laughing Lumberjack Lager | 1           | 14.0000    |
| Outback Lager             | 1           | 15.0000    |
| LakkaLikk                 | 1           | 18.0000    |

# Using the OR operator

- OR requires any one of the conditions to be true

```
SELECT PRODUCT_NAME, UNIT_PRICE, UNITS_IN_STOCK
FROM PRODUCTS
WHERE
UNITS_IN_STOCK = 0
OR
UNIT_PRICE > 100;
```

| PRODUCT_NAME            | UNIT_PRICE | UNITS_IN_STOCK |
|-------------------------|------------|----------------|
| Chef Anton's Gumbo Mix  | 21.3500    | 0              |
| Alice Mutton            | 39.0000    | 0              |
| Thüringer Rostbratwurst | 123.7900   | 0              |
| Gorgonzola Telino       | 12.5000    | 0              |
| Chate de Blaye          | 263.5000   | 17             |
| Perth Pasties           | 32.8000    | 0              |

# Using the OR operator

- NOT is used for complementing the result of the condition

```
SELECT FIRST_NAME, CITY FROM EMPLOYEES
WHERE CITY NOT IN ('SEATTLE', 'LONDON');
```

| FIRST_NAME | CITY     |
|------------|----------|
| Andrew     | Tacoma   |
| Janet      | Kirkland |
| Margaret   | Redmond  |

- You may try NOT BETWEEN, NOT LIKE, IS NOT NULL also.
-

# Precedence control

- Note that the result does not match the expected output due to the default precedence and flow of evaluation (left to right)

```
mysql> select product_name, supplier_id, unit_price
-> from products
-> where unit_price between 10 and 20 and supplier_id=1 or supplier_id=2;
```

| product_name                     | supplier_id | unit_price |
|----------------------------------|-------------|------------|
| Chai                             | 1           | 18.0000    |
| Chang                            | 1           | 19.0000    |
| Aniseed Syrup                    | 1           | 10.0000    |
| Chef Anton's Cajun Seasoning     | 2           | 22.0000    |
| Chef Anton's Gumbo Mix           | 2           | 21.3500    |
| Louisiana Fiery Hot Pepper Sauce | 2           | 21.0500    |
| Louisiana Hot Spiced Okra        | 2           | 17.0000    |

7 rows in set (0.00 sec)

# Precedence control

- Use the parentheses to control the precedence

```
mysql> select product_name, supplier_id, unit_price
-> from products
-> where unit_price between 10 and 20 and (supplier_id=1 or supplier_id=2);
```

| product_name              | supplier_id | unit_price |
|---------------------------|-------------|------------|
| Chai                      | 1           | 18.0000    |
| Chang                     | 1           | 19.0000    |
| Aniseed Syrup             | 1           | 10.0000    |
| Louisiana Hot Spiced Okra | 2           | 17.0000    |

4 rows in set (0.00 sec)

# Sorting the results

- Sort rows with the ORDER BY clause
    - ASC: ascending order, default
    - DESC: descending order
  - The ORDER BY clause comes at the end in the SELECT statement.
-

## Sorting the results

```
SELECT PRODUCT_NAME, CATEGORY_ID, SUPPLIER_ID, UNIT_PRICE
FROM PRODUCTS
WHERE UNIT_PRICE > 40
ORDER BY UNIT_PRICE DESC;
```

| PRODUCT_NAME            | CATEGORY_ID | SUPPLIER_ID | UNIT_PRICE |
|-------------------------|-------------|-------------|------------|
| Chate de Blaye          | 1           | 18          | 263.5000   |
| Thüringer Rostbratwurst | 6           | 12          | 123.7900   |
| Mishi Kobe Niku         | 6           | 4           | 97.0000    |
| Sir Rodney's Marmalade  | 3           | 8           | 81.0000    |
| Carnarvon Tigers        | 8           | 7           | 62.5000    |
| Raclette Courdavault    | 4           | 28          | 55.0000    |
| Manjimup Dried Apples   | 7           | 24          | 53.0000    |
| Tarte au sucre          | 3           | 29          | 49.3000    |
| Ipoh Coffee             | 1           | 20          | 46.0000    |
| Rössle Sauerkraut       | 7           | 12          | 45.6000    |
| Schoggi Schokolade      | 3           | 11          | 43.9000    |
| Vegie-spread            | 2           | 7           | 43.9000    |



# Sorting by Column Alias

```
SELECT PRODUCT_NAME,
UNIT_PRICE * UNITS_IN_STOCK AS STOCK_VALUE
FROM PRODUCTS
ORDER BY STOCK_VALUE DESC LIMIT 10;
```

| PRODUCT_NAME                 | STOCK_VALUE |
|------------------------------|-------------|
| Chate de Blaye               | 4479.5000   |
| Raclette Courdavault         | 4345.0000   |
| Queso Manchego La Pastora    | 3268.0000   |
| Sir Rodney's Marmalade       | 3240.0000   |
| Sirop d'Orangette            | 3220.5000   |
| Grandma's Boysenberry Spread | 3000.0000   |
| Mishi Kobe Niku              | 2813.0000   |
| Pâté chinois                 | 2760.0000   |
| Carnarvon Tigers             | 2625.0000   |
| Boston Crab Meat             | 2263.2000   |

# Sorting by multiple columns

- The second column in the sort-list will be used only if the values in the first column being used for sorting has duplicate values.

```
SELECT PRODUCT_NAME, CATEGORY_ID, UNIT_PRICE
FROM PRODUCTS WHERE UNIT_PRICE > 40
ORDER BY CATEGORY_ID, UNIT_PRICE DESC;
```

| PRODUCT_NAME            | CATEGORY_ID | UNIT_PRICE |
|-------------------------|-------------|------------|
| Chate de Blaye          | 1           | 263.5000   |
| Ipoh Coffee             | 1           | 46.0000    |
| Vegie-spread            | 2           | 43.9000    |
| Sir Rodney's Marmalade  | 3           | 81.0000    |
| Tarte au sucre          | 3           | 49.3000    |
| Schoggi Schokolade      | 3           | 43.9000    |
| Raclette Courdavault    | 4           | 55.0000    |
| Thüringer Rostbratwurst | 6           | 123.7900   |
| Mishi Kobe Niku         | 6           | 97.0000    |
| Manjimup Dried Apples   | 7           | 53.0000    |
| Russle Sauerkraut       | 7           | 45.6000    |
| Carnarvon Tigers        | 8           | 62.5000    |

Note that the data is initially sorted based on category id and then within each category group, the data is sorted based on the price in descending order

# Sorting and Projection

- The column used for sorting need not be in the column list beging retrieved.
  - For example,  
Retrieve the top 5 costliest product names:

```
SELECT PRODUCT_NAME FROM PRODUCTS
ORDER BY UNIT_PRICE DESC
LIMIT 5;
```

| PRODUCT_NAME           |
|------------------------|
| Chte de Blaye          |
| Thringer Rostbratwurst |
| Mishi Kobe Niku        |
| Sir Rodney's Marmalade |
| Carnarvon Tigers       |

# Conditional Expressions – CASE statement

- CASE statement
  - Syntax:
    - CASE value  
    WHEN [compare\_value] THEN result  
    [WHEN [compare\_value] THEN result ...]  
    [ELSE result]  
END
-

# Conditional Expressions – CASE statement

```
SELECT PRODUCT_NAME,
CASE CATEGORY_ID
 WHEN 1 THEN 'BEVERAGES'
 WHEN 2 THEN 'CONDIMENTS'
 WHEN 3 THEN 'CONFECTIONS'
 ELSE 'GENERAL'
END AS CATEGORY
FROM PRODUCTS
ORDER BY CATEGORY;
```

| PRODUCT_NAME                     | CATEGORY    |
|----------------------------------|-------------|
| Chai                             | BEVERAGES   |
| Sasquatch Ale                    | BEVERAGES   |
| Stealthy Stout                   | BEVERAGES   |
| Châte de Blaye                   | BEVERAGES   |
| Ipoh Coffee                      | BEVERAGES   |
| Guaraní Fantástica               | BEVERAGES   |
| Chartreuse verte                 | BEVERAGES   |
| Laughing Lumberjack Lager        | BEVERAGES   |
| Outback Lager                    | BEVERAGES   |
| Rhineland Klosterbier            | BEVERAGES   |
| Chang                            | BEVERAGES   |
| Lakkalikööri                     | BEVERAGES   |
| Aniseed Syrup                    | CONDIMENTS  |
| Chef Anton's Cajun Seasoning     | CONDIMENTS  |
| Chef Anton's Gumbo Mix           | CONDIMENTS  |
| Grandma's Boysenberry Spread     | CONDIMENTS  |
| Gula Malacca                     | CONDIMENTS  |
| Veggie-spread                    | CONDIMENTS  |
| Original Frankfurter green sauce | CONDIMENTS  |
| Sirop d'érable                   | CONDIMENTS  |
| Northwoods Cranberry Sauce       | CONDIMENTS  |
| Louisiana Hot Spiced Okra        | CONDIMENTS  |
| Louisiana Fiery Hot Pepper Sauce | CONDIMENTS  |
| Genen Shoyu                      | CONDIMENTS  |
| Tarte au sucre                   | CONFECTIONS |
| Chocolade                        | CONFECTIONS |
| Zaanse koeken                    | CONFECTIONS |
| Maxilaku                         | CONFECTIONS |
| Scottish Longbreads              | CONFECTIONS |

# Conditional Expressions – IF statement

- IF statement
  - Syntax:
    - IF(expr1,expr2,expr3)
    - If expr1 is TRUE (expr1 <> 0 and expr1 <> NULL)
      - then IF() returns expr2;
      - otherwise it returns expr3.
    - IF() returns a numeric or string value, depending on the context in which it is used.
-

# Conditional Expressions – IF statement

```
SELECT PRODUCT_NAME, UNIT_PRICE,
IF(UNIT_PRICE >= 25, 'COSTLY', 'CHEAP') AS 'IS COSTLY'
FROM PRODUCTS
LIMIT 15;
```

| PRODUCT_NAME                    | UNIT_PRICE | IS COSTLY |
|---------------------------------|------------|-----------|
| Chai                            | 18.0000    | CHEAP     |
| Chang                           | 19.0000    | CHEAP     |
| Aniseed Syrup                   | 10.0000    | CHEAP     |
| Chef Anton's Cajun Seasoning    | 22.0000    | CHEAP     |
| Chef Anton's Gumbo Mix          | 21.3500    | CHEAP     |
| Grandma's Boysenberry Spread    | 25.0000    | COSTLY    |
| Uncle Bob's Organic Dried Pears | 30.0000    | COSTLY    |
| Northwoods Cranberry Sauce      | 40.0000    | COSTLY    |
| Mishi Kobe Niku                 | 97.0000    | COSTLY    |
| Ikura                           | 31.0000    | COSTLY    |
| Queso Cabrales                  | 21.0000    | CHEAP     |
| Queso Manchego La Pastora       | 38.0000    | COSTLY    |
| Konbu                           | 6.0000     | CHEAP     |
| Tofu                            | 23.2500    | CHEAP     |

# Conditional Expressions – IFNULL statement

- IFNULL statement
  - Syntax:
    - IFNULL(expr1,expr2)
    - If expr1 is not NULL,
      - IFNULL() returns expr1;
      - otherwise it returns expr2.
    - IFNULL() returns a numeric or string value, depending on the context in which it is used
-



# Conditional Expressions – IFNULL statement

```
SELECT FIRST_NAME, CITY,
IFNULL(REGION, '--N/A--') AS REGION
FROM EMPLOYEES;
```

| FIRST_NAME | CITY     | REGION  |
|------------|----------|---------|
| Nancy      | Seattle  | WA      |
| Andrew     | Tacoma   | WA      |
| Janet      | Kirkland | WA      |
| Margaret   | Redmond  | WA      |
| Steven     | London   | --N/A-- |
| Michael    | London   | --N/A-- |
| Robert     | London   | --N/A-- |
| Laura      | Seattle  | WA      |
| Anne       | London   | --N/A-- |

# Conditional Expressions – NULLIF statement

- NULLIF statement
  - Syntax:
    - NULLIF(expr1,expr2)
      - Returns NULL if expr1 = expr2 is true,
      - otherwise returns expr1.
-

# Conditional Expressions – NULLIF statement


```
SELECT PRODUCT_NAME, UNIT_PRICE,
NULLIF(UNITS_IN_STOCK, 0) AS UNITS_IN_STOCK FROM PRODUCTS;
```

| PRODUCT_NAME                    | UNIT_PRICE | UNITS_IN_STOCK |
|---------------------------------|------------|----------------|
| Chai                            | 18.0000    | 39             |
| Chang                           | 19.0000    | 17             |
| Aniseed Syrup                   | 10.0000    | 13             |
| Chef Anton's Cajun Seasoning    | 22.0000    | 53             |
| Chef Anton's Gumbo Mix          | 21.3500    | NULL           |
| Grandma's Boysenberry Spread    | 25.0000    | 120            |
| Uncle Bob's Organic Dried Pears | 30.0000    | 15             |
| Northwoods Cranberry Sauce      | 40.0000    | 6              |
| Mishi Kobe Niku                 | 97.0000    | 29             |
| Ikura                           | 31.0000    | 31             |
| Queso Cabrales                  | 21.0000    | 22             |
| Queso Manchego La Pastora       | 38.0000    | 86             |
| Konbu                           | 6.0000     | 24             |
| Tofu                            | 23.2500    | 35             |
| Genen Shouyu                    | 15.5000    | 39             |
| Pavlova                         | 17.4500    | 29             |
| Alice Mutton                    | 39.0000    | NULL           |
| Carnarvon Tigers                | 62.5000    | 42             |

# Functions

- A function is a callable unit of code to perform a task
  - A function may take 0 or more arguments (parameters), which can be fed as input to the job to be done
  - A function returns a value, which usually represents the result of the job being done
  - For example,
    - The sysdate() function does not take any parameter, but returns the current system date and time:  
`SELECT SYSDATE();`
    - The length() function takes one string parameter and returns the no.of characters in it:  
`SELECT LENGTH('MYSQL IS FUN');`
-

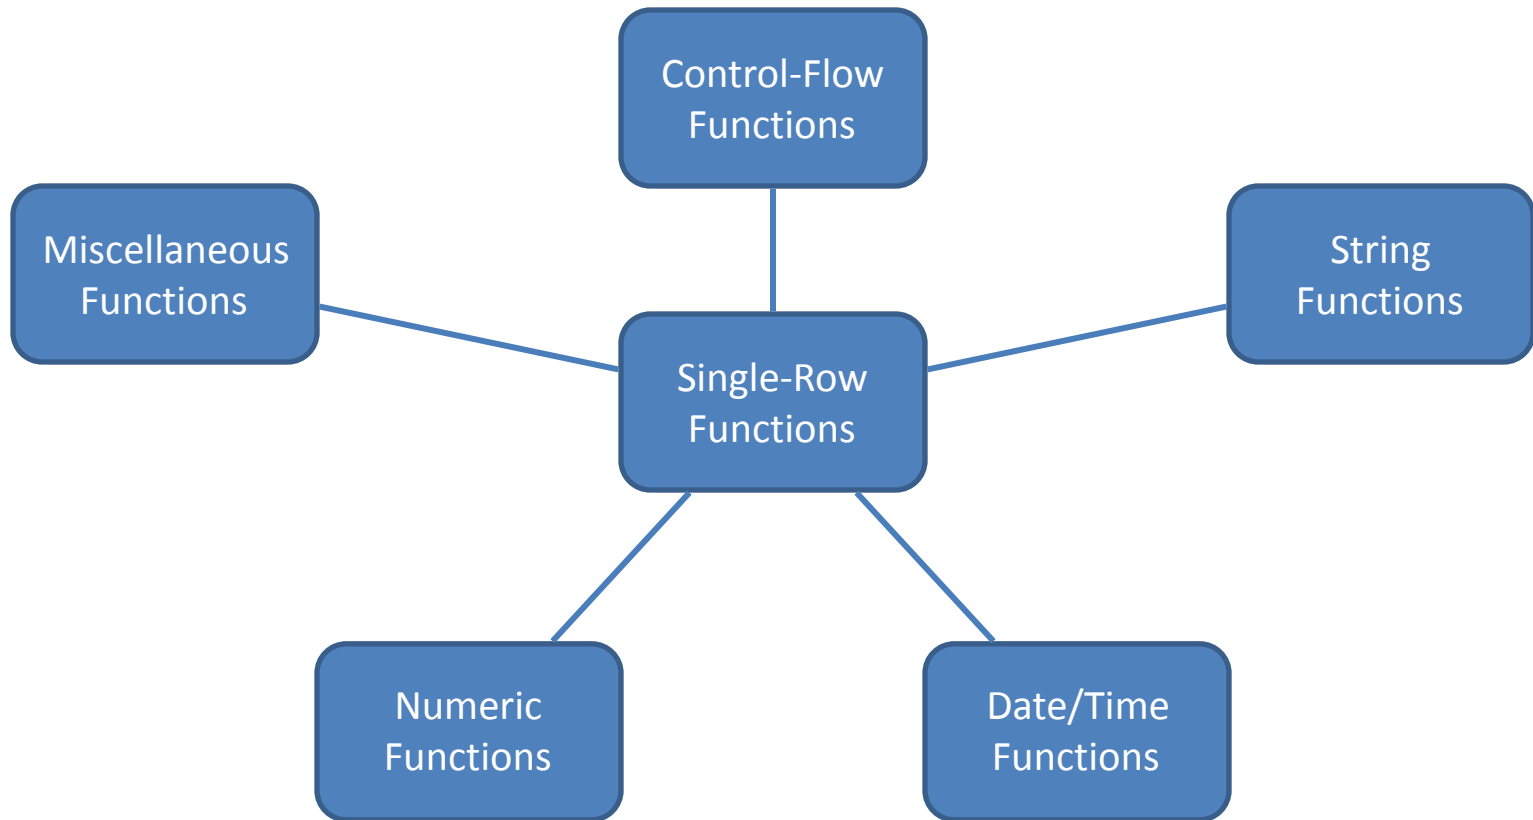
# Two types of functions



# Single-Row functions

- Single row functions:
    - Manipulate data items
    - Accept arguments and return one value
    - Act on each row returned
    - Return one result per row
    - May modify the data type
    - Can be nested
    - Accept arguments which can be a column or an expression
-

# Single-Row functions



# Control-Flow functions

- IF()
  - Used for If/Else construct
- IFNULL()
  - Null-if/else construct
  - Used for NULL value substitution
- NULLIF()
  - Return NULL if `expr1 = expr2`
- This topic has already been discussed in the last session.



# String functions

| Function                            | Description                                                                                                        |
|-------------------------------------|--------------------------------------------------------------------------------------------------------------------|
| CHAR()                              | Return the character for each integer passed                                                                       |
| CHAR_LENGTH()<br>CHARACTER_LENGTH() | Returns the length of the string str, measured in characters. A multi-byte character counts as a single character. |
| CONCAT()<br>CONCAT_WS()             | CONCAT_WS() stands for Concatenate With Separator and is a special form of CONCAT().                               |
| FORMAT()                            | Return a number formatted to specified number of decimal places                                                    |
| LOWER()<br>LCASE()                  | Return the argument in lowercase                                                                                   |
| LEFT()<br>RIGHT()                   | Return the leftmost (rightmost) number of characters as specified                                                  |
| LENGTH()                            | Return the length of a string in bytes                                                                             |
| LOAD()                              | Load the named file. Used for inserting into BLOB fields.                                                          |
| LOCATE()<br>POSITION()              | Return the position of the first occurrence of substring                                                           |
| LPAD()<br>RPAD()                    | Return the string argument, left (or right) padded with the specified string                                       |

# String functions

| Function                     | Description                                       |
|------------------------------|---------------------------------------------------|
| LTRIM()<br>RTRIM()<br>TRIM() | Removes extra spaces (left, right or both)        |
| SPACE()                      | Return a string of the specified number of spaces |
| SUBSTR()                     | Return the substring as specified                 |
| UPPER()<br>UCASE()           | Convert to uppercase                              |

# String functions - Examples

```
SELECT UPPER(CUSTOMER_NAME) FROM CUSTOMERS LIMIT 10;
SELECT LOWER(CUSTOMER_NAME) FROM CUSTOMERS LIMIT 10;
```

```
-- Case insensitive comparison
SELECT * FROM CUSTOMERS
 WHERE UPPER(CITY) = 'LONDON';
```

```
SELECT
CONCAT(FIRST_NAME, LAST_NAME, 'lives in', CITY) AS INFO1,
CONCAT_WS(' ', FIRST_NAME, LAST_NAME, 'lives in', CITY) AS INFO2
FROM EMPLOYEES;
```

| INFO1                          | INFO2                             |
|--------------------------------|-----------------------------------|
| NancyDavoliolives inSeattle    | Nancy Davolio lives in Seattle    |
| AndrewFullerlives inTacoma     | Andrew Fuller lives in Tacoma     |
| JanetLeverlinglives inKirkland | Janet Leverling lives in Kirkland |
| MargaretPeacocklives inRedmond | Margaret Peacock lives in Redmond |
| StevenBuchananlives inLondon   | Steven Buchanan lives in London   |
| MichaelSuyamalives inLondon    | Michael Suyama lives in London    |
| RobertKinglives inLondon       | Robert King lives in London       |
| LauraCallahanlives inSeattle   | Laura Callahan lives in Seattle   |
| AnneDodsworthlives inLondon    | Anne Dodsworth lives in London    |

# Date/time functions

| Function                                                                                                                           | Description                                  |
|------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------|
| CURDATE()<br>CURRENT_DATE()<br>CURRENT_DATE                                                                                        | Return the current date                      |
| CURTIME()<br>CURRENT_TIME()<br>CURRENT_TIME                                                                                        | Return the current time                      |
| CURRENT_TIMESTAMP()<br>CURRENT_TIMESTAMP()<br>NOW()<br>SYSDATE()<br>LOCALTIME()<br>LOCALTIME<br>LOCALTIMESTAMP()<br>LOCALTIMESTAMP | Return the current date and time             |
| DATE_ADD()                                                                                                                         | Add time values (intervals) to a date value  |
| DATE_SUB()                                                                                                                         | Subtract a time value (interval) from a date |
| DATEDIFF()                                                                                                                         | Subtract two dates                           |
| MAKEDATE()                                                                                                                         | Create a date from the year and day of year  |

# Date/time functions

| Function     | Description                              |
|--------------|------------------------------------------|
| DAYNAME()    | Return the name of the weekday           |
| DAYOFMONTH() | Return the day of the month (0-31)       |
| DAYOFWEEK()  | Return the weekday index of the argument |
| DAYOFYEAR()  | Return the day of the year (1-366)       |

# Date/Time functions - Examples

```
SELECT SYSDATE(), DAYOFYEAR(SYSDATE()), DAYNAME(SYSDATE());
```

| SYSDATE()           | DAYOFYEAR(SYSDATE()) | DAYNAME(SYSDATE()) |
|---------------------|----------------------|--------------------|
| 2012-07-18 17:27:13 | 200                  | Wednesday          |

```
SELECT MAKEDATE(2012, 200);
```

| MAKEDATE(2012, 200) |
|---------------------|
| 2012-07-18          |

# Date/Time functions - Examples

```
SELECT DATE_ADD('2012-01-01', INTERVAL 365 DAY);
```

| DATE_ADD('2012-01-01', INTERVAL 365 DAY) |
|------------------------------------------|
| 2012-12-31                               |

```
SELECT NOW(), DATE_ADD(NOW(), INTERVAL 10 HOUR);
```

| NOW()               | DATE_ADD(NOW(), INTERVAL 10 HOUR) |
|---------------------|-----------------------------------|
| 2012-07-18 17:31:22 | 2012-07-19 03:31:22               |

For more information about INTERVAL,

Visit [http://dev.mysql.com/doc/refman/5.1/en/date-and-time-functions.html#function\\_date-add](http://dev.mysql.com/doc/refman/5.1/en/date-and-time-functions.html#function_date-add)

---

# Date/Time functions - Examples

```
SELECT
CONCAT_WS(' ', FIRST_NAME, LAST_NAME, 'has',
 DATEDIFF(SYSDATE(), HIRE_DATE) DIV 365,
 'years experience') AS INFO
FROM EMPLOYEES;
```

```
+-----+
| INFO |
+-----+
| Nancy Davolio has 20 years experience |
| Andrew Fuller has 19 years experience |
| Janet Leverling has 20 years experience |
| Margaret Peacock has 19 years experience|
| Steven Buchanan has 18 years experience |
| Michael Suyama has 18 years experience |
| Robert King has 18 years experience |
| Laura Callahan has 18 years experience |
| Anne Dodsworth has 17 years experience |
+-----+
```



# Numeric functions

| Function       | Description                                                    |
|----------------|----------------------------------------------------------------|
| ABS()          | Return the absolute value                                      |
| ROUND()        | Round the argument                                             |
| RAND()         | Return a random floating-point value                           |
| POW(), POWER() | Return the argument raised to the specified power              |
| SIGN()         | Return the sign of the argument                                |
| TRUNCATE()     | Truncate to specified number of decimal places                 |
| CEIL()         | Return the smallest integer value not less than the argument   |
| FLOOR()        | Return the largest integer value not greater than the argument |
| MOD()          | Return the remainder                                           |
| CONV()         | Convert numbers between different number bases                 |

Complete reference: <http://dev.mysql.com/doc/refman/5.1/en/mathematical-functions.html>

---

# Numeric functions - Examples

```
SELECT CEIL(2.16), CEIL(2.96),
 FLOOR(2.16), FLOOR(2.96);
```

| CEIL(2.16) | CEIL(2.96) | FLOOR(2.16) | FLOOR(2.96) |
|------------|------------|-------------|-------------|
| 3          | 3          | 2           | 2           |

```
SELECT ROUND(2.16, 1), ROUND(2.96, 1),
 TRUNCATE(2.16, 1), TRUNCATE(2.96, 1);
```

| ROUND(2.16, 1) | ROUND(2.96, 1) | TRUNCATE(2.16, 1) | TRUNCATE(2.96, 1) |
|----------------|----------------|-------------------|-------------------|
| 2.2            | 3.0            | 2.1               | 2.9               |

```
SELECT CONV(1974, 10, 2), CONV(1974, 10, 8), CONV(1974, 10, 16);
```

| CONV(1974, 10, 2) | CONV(1974, 10, 8) | CONV(1974, 10, 16) |
|-------------------|-------------------|--------------------|
| 11110110110       | 3666              | 7B6                |

```
SELECT CONV(11110110110, 2, 10), CONV(3666, 8, 10), CONV('7B6', 16, 10);
```

| CONV(11110110110, 2, 10) | CONV(3666, 8, 10) | CONV('7B6', 16, 10) |
|--------------------------|-------------------|---------------------|
| 1974                     | 1974              | 1974                |

# Numeric functions - Examples

```
SELECT 12 MOD 5, MOD(12, 5);
```

| 12 MOD 5 | MOD(12, 5) |
|----------|------------|
| 2        | 2          |

```
SELECT RAND(), RAND(), RAND();
```

| RAND()             | RAND()            | RAND()              |
|--------------------|-------------------|---------------------|
| 0.6117338664878088 | 0.653780472464506 | 0.43330319359274877 |

```
SELECT POW(2, 10), POWER(2, 10);
```

| POW(2, 10) | POWER(2, 10) |
|------------|--------------|
| 1024       | 1024         |

---

# Miscellaneous functions

| Function                                                                    | Description                                              |
|-----------------------------------------------------------------------------|----------------------------------------------------------|
| CURRENT_USER()<br>CURRENT_USER<br>SESSION_USER()<br>SYSTEM_USER()<br>USER() | The user name and host name provided by the client       |
| DATABASE()<br>SCHEMA()                                                      | Return the default (current) database name               |
| VERSION()                                                                   | Returns a string that indicates the MySQL server version |
| DEFAULT()                                                                   | Return the default value for a table column              |
| ENCODE(), DECODE()                                                          | Encode/Decode a string                                   |
| MD5()                                                                       | Calculate MD5 checksum                                   |
| PASSWORD()                                                                  | Calculate and return a password string (one-way)         |

Complete reference: <http://dev.mysql.com/doc/refman/5.1/en/mathematical-functions.html>

---

# Miscellaneous functions - Examples

```
SELECT USER(), DATABASE(), VERSION();
```

| USER()         | DATABASE() | VERSION() |
|----------------|------------|-----------|
| root@localhost | _prep      | 5.5.15    |

```
SELECT * FROM PERSONS;
```

| ID | NAME   | CITY      |
|----|--------|-----------|
| 1  | VINOD  | BANGALORE |
| 2  | SCOTT  | DALLAS    |
| 3  | BANU   | BANGALORE |
| 4  | MILLER | CHICAGO   |

```
SELECT * FROM PERSONS
WHERE CITY = DEFAULT(CITY);
```

| ID | NAME  | CITY      |
|----|-------|-----------|
| 1  | VINOD | BANGALORE |
| 3  | BANU  | BANGALORE |

```
SELECT * FROM USERS;
```

| USERNAME | PASSWD                                    |
|----------|-------------------------------------------|
| VINOD    | *9D2D1415F6A8D138D693228417BA3ED2B7D06C9E |
| SCOTT    | *746FC1265EFA3E5BC598A1F2122C3377D74E933D |
| ROOT     | *B481E6BC13C8A73BC4C2B05F6F1393D16F2EE840 |

```
SELECT USERNAME FROM USERS
WHERE PASSWD = PASSWORD('SECRET');
```

| USERNAME |
|----------|
| VINOD    |

Aggregating data

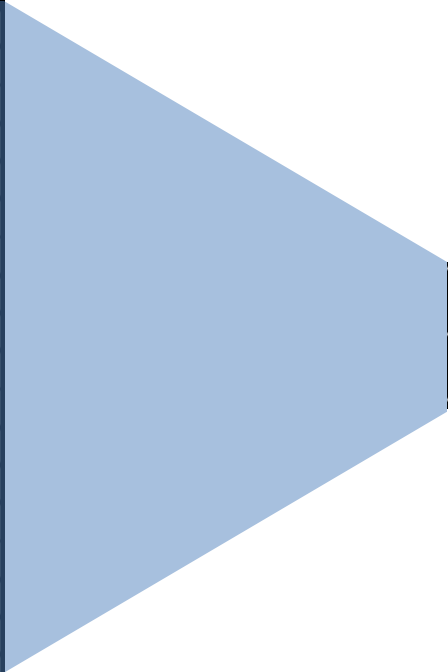
# Using Group Functions

---

# Group Functions

- Group functions operate on sets of rows to give one result per group.

| PRODUCT_NAME                    | UNIT_PRICE |
|---------------------------------|------------|
| Chai                            | 18.0000    |
| Chang                           | 19.0000    |
| Aniseed Syrup                   | 10.0000    |
| Chef Anton's Cajun Seasoning    | 22.0000    |
| Chef Anton's Gumbo Mix          | 21.3500    |
| Grandma's Boysenberry Spread    | 25.0000    |
| Uncle Bob's Organic Dried Pears | 30.0000    |
| Northwoods Cranberry Sauce      | 40.0000    |
| Mishi Kobe Niku                 | 97.0000    |
| Ikura                           | 31.0000    |
| Queso Cabrales                  | 21.0000    |
| Queso Manchego La Pastora       | 38.0000    |
| Konbu                           | 6.0000     |
| Tofu                            | 23.2500    |
| Genen Shouyu                    | 15.5000    |
| Pavlova                         | 17.4500    |
| Alice Mutton                    | 39.0000    |
| Corned Beef                     | 26.5000    |



|                 |
|-----------------|
| AVG(UNIT_PRICE) |
| 28.86636364     |

# Group Functions

- Some of the most commonly used group functions:
    - AVG
    - COUNT
    - MAX
    - MIN
    - SUM
-



# Group Functions

```
SELECT [column,] group_function(column), ...
FROM table
[WHERE condition]
[GROUP BY column [HAVING {condition}]]
[ORDER BY column];
```

---

# Group Functions

```
SELECT MAX(UNIT_PRICE), AVG(UNIT_PRICE), MIN(UNIT_PRICE), MAX(UNIT_PRICE)
FROM PRODUCTS;
```

| MAX(UNIT_PRICE) | AVG(UNIT_PRICE) | MIN(UNIT_PRICE) | MAX(UNIT_PRICE) |
|-----------------|-----------------|-----------------|-----------------|
| 263.5000        | 28.86636364     | 2.5000          | 263.5000        |

You can use AVG and SUM for numeric data only, where as, MIN and MAX can be used even on any data type.

```
SELECT MIN(HIRE_DATE), MAX(HIRE_DATE) FROM EMPLOYEES;
```

| MIN(HIRE_DATE)      | MAX(HIRE_DATE)      |
|---------------------|---------------------|
| 1992-04-01 00:00:00 | 1994-11-15 00:00:00 |

# Group Functions

When used on strings, MIN and MAX compare the ASCII value of the characters

```
SELECT MIN(FIRST_NAME), MAX(FIRST_NAME)
FROM EMPLOYEES;
```

| MIN(FIRST_NAME) | MAX(FIRST_NAME) |
|-----------------|-----------------|
| Andrew          | Steven          |

# Group Functions

- COUNT(\*) returns the number of rows in a table.
- COUNT(expr) returns the number of rows with non-null values for the expr.
- COUNT(DISTINCT expr) returns the number of distinct non-null values of the expr.

```
SELECT REGION
FROM EMPLOYEES;
```

| REGION |
|--------|
| WA     |
| WA     |
| WA     |
| WA     |
| NULL   |
| NULL   |
| NULL   |
| WA     |
| NULL   |

```
SELECT COUNT(*), COUNT(REGION),
COUNT(DISTINCT REGION) FROM EMPLOYEES;
```

| COUNT(*) | COUNT(REGION) | COUNT(DISTINCT REGION) |
|----------|---------------|------------------------|
| 9        | 5             | 1                      |

# Group Functions

- Group functions ignore null values in the column.
- For example, the following statement would ignore any NULL values in the UNIT\_PRICE column, and perform the calculation

```
SELECT AVG(UNIT_PRICE) FROM PRODUCTS
```

- It is same as:  
$$\text{SUM\_OF\_NON\_NULL\_VALUES} / \text{COUNT\_OF\_NON\_NULL\_VALUES}$$

# Group Functions

- If you don't want to ignore the NULL values, you may use the IFNULL() function as shown below:

```
SELECT AVG(IFNULL(UNIT_PRICE, 0)) FROM PRODUCTS;
```

# Creating Groups Of Data

```
SELECT FIRST_NAME, LAST_NAME, CITY
FROM EMPLOYEES ORDER BY CITY;
```

```
SELECT CITY, COUNT(*)
FROM EMPLOYEES
GROUP BY CITY;
```

| FIRST_NAME | LAST_NAME | CITY     |
|------------|-----------|----------|
| Janet      | Leverling | Kirkland |
| Steven     | Buchanan  | London   |
| Michael    | Suyama    | London   |
| Robert     | King      | London   |
| Anne       | Dodsworth | London   |
| Margaret   | Peacock   | Redmond  |
| Nancy      | Davolio   | Seattle  |
| Laura      | Callahan  | Seattle  |
| Andrew     | Fuller    | Tacoma   |

All employees

| CITY     | COUNT(*) |
|----------|----------|
| Kirkland | 1        |
| London   | 4        |
| Redmond  | 1        |
| Seattle  | 2        |
| Tacoma   | 1        |

The number of employees in each city

# Creating Groups Of Data

- Divide rows in a table into smaller groups by using the GROUP BY clause.

```
SELECT column, group_function(column), ...
FROM table
[WHERE condition]
GROUP BY column [HAVING {condition}]]
[ORDER BY column],
```

---



# Creating Groups Of Data

- All columns in the SELECT list that are not group functions must be in the GROUP BY clause.

```
SELECT FIRST_NAME, CITY, COUNT(*)
FROM EMPLOYEES GROUP BY CITY;
```

| FIRST_NAME | CITY     | COUNT(*) |
|------------|----------|----------|
| Janet      | Kirkland | 1        |
| Steven     | London   | 4        |
| Margaret   | Redmond  | 1        |
| Nancy      | Seattle  | 2        |
| Andrew     | Tacoma   | 1        |

Invalid query. Unlike other DBMS, MySQL does not give error for this. Instead, it just picks the first employee name in each group and outputs.

```
SELECT CITY, COUNT(*)
FROM EMPLOYEES GROUP BY CITY;
```

| CITY     | COUNT(*) |
|----------|----------|
| Kirkland | 1        |
| London   | 4        |
| Redmond  | 1        |
| Seattle  | 2        |
| Tacoma   | 1        |

Valid query.

# Creating Groups Of Data

- The GROUP BY column does not have to be in the SELECT list.

```
SELECT COUNT(*) FROM PRODUCTS
GROUP BY CATEGORY_ID;
```

| COUNT(*) |
|----------|
| 12       |
| 12       |
| 13       |
| 10       |
| 7        |
| 6        |
| 5        |
| 12       |

- However, such queries are very rarely used.
-

# Grouping by More Than One Column

- Example – Total stock value of all products of a particular category, supplied by a particular supplier

```
SELECT CATEGORY_ID, SUPPLIER_ID, SUM(UNIT_PRICE * UNITS_IN_STOCK) AS INFO
FROM PRODUCTS GROUP BY CATEGORY_ID, SUPPLIER_ID;
```

| NAME                   | CATEGORY_ID | SUPPLIER_ID | UNIT_PRICE | UNITS_IN_STOCK |
|------------------------|-------------|-------------|------------|----------------|
|                        | 1           | 1           | 18.0000    | 39             |
|                        | 1           | 1           | 19.0000    | 17             |
| lager                  | 1           | 7           | 15.0000    | 15             |
| ant-istica             | 1           | 10          | 4.5000     | 20             |
| Klosterbier            | 1           | 12          | 7.7500     | 125            |
| tout                   | 1           | 16          | 18.0000    | 20             |
| Ale                    | 1           | 16          | 14.0000    | 111            |
| umberjack Lager        | 1           | 16          | 14.0000    | 52             |
| e verte                | 1           | 18          | 18.0000    | 69             |
| laye                   | 1           | 18          | 263.5000   | 17             |
| e                      | 1           | 20          | 46.0000    | 17             |
| Hri                    | 1           | 23          | 18.0000    | 57             |
| rup                    | 2           | 1           | 10.0000    | 13             |
| Fiery Hot Pepper Sauce | 2           | 2           | 21.0000    | 76             |
| Hot Spiced Okra        | 2           | 2           | 17.0000    | 4              |
| 's Cajun Seasoning     | 2           | 2           | 22.0000    | 53             |
| 's Gumbo Mix           | 2           | 2           | 21.3500    | 0              |
| Boysenberry Spread     | 2           | 3           | 25.0000    | 120            |

| CATEGORY_ID | SUPPLIER_ID | INFO      |
|-------------|-------------|-----------|
| 1           | 1           | 1025.0000 |
| 1           | 7           | 225.0000  |
| 1           | 10          | 90.0000   |
| 1           | 12          | 968.7500  |
| 1           | 16          | 2642.0000 |
| 1           | 18          | 5721.5000 |
| 1           | 20          | 782.0000  |
| 1           | 23          | 1026.0000 |
| 2           | 1           | 130.0000  |
| 2           | 2           | 2833.8000 |
| 2           | 3           | 3240.0000 |
| 2           | 6           | 604.5000  |
| 2           | 7           | 1052.6000 |

# Restricting the groups using HAVING clause

- You cannot use the WHERE clause to restrict groups.
  - You cannot use group functions in the WHERE clause.

```
mysql> SELECT CATEGORY_ID, SUPPLIER_ID, SUM(UNIT_PRICE * UNITS_IN_STOCK)
-> FROM PRODUCTS
-> WHERE SUM(UNIT_PRICE*UNITS_IN_STOCK) >=2500
-> GROUP BY CATEGORY_ID, SUPPLIER_ID;
ERROR 1111 (HY000): Invalid use of group function
```

# Restricting the groups using HAVING clause

- You use the HAVING clause to restrict groups.

```
SELECT [column,] group_function(column), ...
FROM table
[WHERE condition]
[GROUP BY column [HAVING {condition}]]
[ORDER BY column];
```

```
SELECT CATEGORY_ID, SUPPLIER_ID, SUM(UNIT_PRICE * UNITS_IN_STOCK)
FROM PRODUCTS GROUP BY CATEGORY_ID, SUPPLIER_ID
HAVING SUM(UNIT_PRICE*UNITS_IN_STOCK) >=2500;
```

| CATEGORY_ID | SUPPLIER_ID | SUM(UNIT_PRICE * UNITS_IN_STOCK) |
|-------------|-------------|----------------------------------|
| 1           | 16          | 2642.0000                        |
| 1           | 18          | 5721.5000                        |
| 2           | 2           | 2833.8000                        |
| 2           | 3           | 3240.0000                        |
| 2           | 29          | 3220.5000                        |
| 3           | 8           | 3575.0000                        |
| 3           | 11          | 3683.5500                        |
| 4           | 5           | 3730.0000                        |
| 4           | 28          | 4991.0000                        |
| 5           | 9           | 2733.0000                        |
| 6           | 4           | 2813.0000                        |
| 5           | 75          | 7815.4500                        |

# Restricting the groups using HAVING clause

- You can use the alias in the HAVING clause
  - This may not work in other RDBMS

```
SELECT CATEGORY_ID, SUPPLIER_ID, SUM(UNIT_PRICE * UNITS_IN_STOCK) AS INFO
FROM PRODUCTS
GROUP BY CATEGORY_ID, SUPPLIER_ID
HAVING INFO >= 2500;
```

| CATEGORY_ID | SUPPLIER_ID | INFO      |
|-------------|-------------|-----------|
| 1           | 16          | 2642.0000 |
| 1           | 18          | 5721.5000 |
| 2           | 2           | 2833.8000 |
| 2           | 3           | 3240.0000 |
| 2           | 29          | 3220.5000 |
| 3           | 8           | 3575.0000 |
| 3           | 11          | 3683.5500 |
| 4           | 5           | 3730.0000 |
| 4           | 28          | 4991.0000 |
| 5           | 9           | 2733.0000 |
| 6           | 4           | 2813.0000 |
| 6           | 25          | 2916.4500 |

Displaying Data From Multiple Tables

# SQL JOINS



## PRORUCTS

| PRODUCT_ID | PRODUCT_NAME                    | CATEGORY_ID |
|------------|---------------------------------|-------------|
| 1          | Chai                            | 1           |
| 2          | Chang                           | 1           |
| 3          | Aniseed Syrup                   | 2           |
| 4          | Chef Anton's Cajun Seasoning    | 2           |
| 5          | Chef Anton's Gumbo Mix          | 2           |
| 6          | Grandma's Boysenberry Spread    | 2           |
| 7          | Uncle Bob's Organic Dried Pears | 7           |
| 8          | Northwoods Cranberry Sauce      | 2           |
| 9          | Mishi Kobe Niku                 | 6           |
| 10         | Ikura                           | 8           |
| 11         | Queso Cabrales                  | 4           |

## CATEGORIES

| CATEGORY_ID | CATEGORY_NAME  | DESCRIPTION |
|-------------|----------------|-------------|
| 1           | Beverages      | Soft drink  |
| 2           | Condiments     | Sweet and   |
| 3           | Confections    | Desserts,   |
| 4           | Dairy Products | Cheeses     |
| 5           | Grains/Cereals | Breads, cr  |
| 6           | Meat/Poultry   | Prepared m  |
| 7           | Produce        | Dried fruit |
| 8           | Seafood        | Seaweed an  |

## PRODUCTS JOINED WITH CATEGORIES

| PRODUCT_ID | PRODUCT_NAME                    | CATEGORY_NAME  |
|------------|---------------------------------|----------------|
| 1          | Chai                            | Beverages      |
| 2          | Chang                           | Beverages      |
| 3          | Aniseed Syrup                   | Condiments     |
| 4          | Chef Anton's Cajun Seasoning    | Condiments     |
| 5          | Chef Anton's Gumbo Mix          | Condiments     |
| 6          | Grandma's Boysenberry Spread    | Condiments     |
| 7          | Uncle Bob's Organic Dried Pears | Produce        |
| 8          | Northwoods Cranberry Sauce      | Condiments     |
| 9          | Mishi Kobe Niku                 | Meat/Poultry   |
| 10         | Ikura                           | Seafood        |
| 11         | Queso Cabrales                  | Dairy Products |



# Cartesian Product

- A Cartesian product is formed when:
    - A join condition is omitted
    - A join condition is invalid
  - All rows in the first table are joined to all rows in the second table
  - To avoid a Cartesian product, always include a valid join condition in a WHERE clause.
-

# Cartesian Product

```
mysql> SELECT COUNT(*) FROM PRODUCTS;
+-----+
| COUNT(*) |
+-----+
| 77 |
+-----+
1 row in set (0.00 sec)

mysql> SELECT COUNT(*) FROM CATEGORIES;
+-----+
| COUNT(*) |
+-----+
| 8 |
+-----+
1 row in set (0.00 sec)

mysql> SELECT COUNT(*) FROM PRODUCTS, CATEGORIES;
+-----+
| COUNT(*) |
+-----+
| 616 |
+-----+
1 row in set (0.00 sec)
```

77 Products X 8 Categories

= 616 Records

# Cartesian Product

```
SELECT PRODUCT_NAME, CATEGORY_NAME
FROM PRODUCTS, CATEGORIES;
```

As you can see, each product is displayed with each categories, since there is no mapping between the product and the category it belongs to.

| PRODUCT_NAME                    | CATEGORY_NAME  |
|---------------------------------|----------------|
| Chai                            | Beverages      |
| Chai                            | Condiments     |
| Chai                            | Confections    |
| Chai                            | Dairy Products |
| Chai                            | Grains/Cereals |
| Chai                            | Meat/Poultry   |
| Chai                            | Produce        |
| Chai                            | Seafood        |
| Chang                           | Beverages      |
| Chang                           | Condiments     |
| Chang                           | Confections    |
| Chang                           | Dairy Products |
| Chang                           | Grains/Cereals |
| Chang                           | Meat/Poultry   |
| Chang                           | Produce        |
| Chang                           | Seafood        |
| Rhönbräu Klosterbier            | Seafood        |
| Lakkalik                        | Beverages      |
| Lakkalik                        | Condiments     |
| Lakkalik                        | Confections    |
| Lakkalik                        | Dairy Products |
| Lakkalik                        | Grains/Cereals |
| Lakkalik                        | Meat/Poultry   |
| Lakkalik                        | Produce        |
| Lakkalik                        | Seafood        |
| Original Frankfurter grüne Soße | Beverages      |
| Original Frankfurter grüne Soße | Condiments     |
| Original Frankfurter grüne Soße | Confections    |
| Original Frankfurter grüne Soße | Dairy Products |
| Original Frankfurter grüne Soße | Grains/Cereals |
| Original Frankfurter grüne Soße | Meat/Poultry   |
| Original Frankfurter grüne Soße | Produce        |
| Original Frankfurter grüne Soße | Seafood        |

516 rows in set (0.00 sec)

# Avoiding Cartesian Product

- Use a join to query data from more than one table.

```
SELECT table1.column, table2.column
FROM table1, table2
WHERE table1.column1 = table2.column2;
```

- Write the join condition in the WHERE clause.
- Prefix the column name with the table name when the same column name appears in more than one table.

```
SELECT PRODUCT_ID, PRODUCT_NAME, CATEGORY_NAME
FROM PRODUCTS, CATEGORIES
WHERE PRODUCTS.CATEGORY_ID = CATEGORIES.CATEGORY_ID;
```

```
SELECT PRODUCT_ID, PRODUCT_NAME, CATEGORY_NAME
FROM PRODUCTS AS P, CATEGORIES AS C
WHERE P.CATEGORY_ID = C.CATEGORY_ID;
```

---

# What is Equi-Join?

- Foreign key of child table is compared with primary key of parent table using the “=” operator
- WHERE PRODUCTS.CATEGORY\_ID = CATEGORIES.CATEGORY\_ID

PRODUCTS – CHILD TABLE

| PRODUCT_ID | PRODUCT_NAME                    | CATEGORY_ID |
|------------|---------------------------------|-------------|
| 1          | Chai                            | 1           |
| 2          | Chang                           | 1           |
| 3          | Aniseed Syrup                   | 2           |
| 4          | Chef Anton's Cajun Seasoning    | 2           |
| 5          | Chef Anton's Gumbo Mix          | 2           |
| 6          | Grandma's Boysenberry Spread    | 2           |
| 7          | Uncle Bob's Organic Dried Pears | 7           |
| 8          | Northwoods Cranberry Sauce      | 2           |
| 9          | Mishi Kobe Niku                 | 6           |
| 10         | Ikura                           | 8           |
| 11         | Queso Cabrales                  | 4           |

Foreign  
key

CATEGORIES – PARENT TABLE

| CATEGORY_ID | CATEGORY_NAME  | DESCRIPTION |
|-------------|----------------|-------------|
| 1           | Beverages      | Soft drink  |
| 2           | Condiments     | Sweet and   |
| 3           | Confections    | Desserts,   |
| 4           | Dairy Products | Cheeses     |
| 5           | Grains/Cereals | Breads, cr  |
| 6           | Meat/Poultry   | Prepared m  |
| 7           | Produce        | Dried fruit |
| 8           | Seafood        | Seaweed an  |

Primary  
key

# Using the ANSI JOIN syntax

- ANSI standardizes the way a join has to be performed using the following syntax:

```
SELECT column_list
FROM table_1 JOIN table_2 ON join_condition
```

- Use the WHERE clause only to filter the rows and not to do the JOIN operation.

```
SELECT column_list
FROM table_1 JOIN table_2 ON join_condition
WHERE filter_condition
```

---

- Non-Standard way of joining:

```
SELECT PRODUCT_ID, PRODUCT_NAME, CATEGORY_NAME
FROM PRODUCTS AS P, CATEGORIES AS C
WHERE P.CATEGORY_ID = C.CATEGORY_ID;
```

- ANSI style (standard):

```
SELECT PRODUCT_ID, PRODUCT_NAME, CATEGORY_NAME
FROM PRODUCTS AS P JOIN CATEGORIES AS C
ON P.CATEGORY_ID = C.CATEGORY_ID;
```

---

# Joining more than two tables

- You can join any number of tables using the following syntax:

```
SELECT column_list
FROM table_1 JOIN table_2 ON join_condition
JOIN table_3 ON join_condition
JOIN table_3 ON join_condition
JOIN table_3 ON join_condition
...
```

---



# Joining more than two tables - Example

- Retrieve the name of the product along with its category name and supplier name for all products

```
SELECT PRODUCT_NAME, CATEGORY_NAME, COMPANY_NAME
FROM PRODUCTS AS P JOIN CATEGORIES AS C ON P.CATEGORY_ID = C.CATEGORY_ID
JOIN SUPPLIERS AS S ON P.SUPPLIER_ID = S.SUPPLIER_ID;
```

| PRODUCT_NAME                 | CATEGORY_NAME | COMPANY_NAME                      |
|------------------------------|---------------|-----------------------------------|
| Chai                         | Beverages     | Exotic Liquids                    |
| Chang                        | Beverages     | Exotic Liquids                    |
| Guaraní Fantástica           | Beverages     | Refrescos Americanas LTDA         |
| Sasquatch Ale                | Beverages     | Bigfoot Breweries                 |
| Steeleye Stout               | Beverages     | Bigfoot Breweries                 |
| Chte de Blaye                | Beverages     | Aux joyeux ecclésiastiques        |
| Chartreuse verte             | Beverages     | Aux joyeux ecclésiastiques        |
| Ipoh Coffee                  | Beverages     | Leka Trading                      |
| Laughing Lumberjack Lager    | Beverages     | Bigfoot Breweries                 |
| Outback Lager                | Beverages     | Pavlova, Ltd.                     |
| Rheinbräu Klosterbier        | Beverages     | Plutzer Lebensmittelgroßmärkte AG |
| Lakkalikööri                 | Beverages     | Karkki Oy                         |
| Aniseed Syrup                | Condiments    | Exotic Liquids                    |
| Chef Anton's Cajun Seasoning | Condiments    | New Orleans Cajun Delights        |
| Chef Anton's Gumbo Mix       | Condiments    | New Orleans Cajun Delights        |
| Grandma's Boysenberry Spread | Condiments    | Grandma Kelly's Homestead         |

# Joining more than two tables - Example

- Retrieve the name of the product along with its category name and supplier name for all products priced above \$50

```
SELECT PRODUCT_NAME, CATEGORY_NAME, COMPANY_NAME, UNIT_PRICE
FROM PRODUCTS AS P JOIN CATEGORIES AS C ON P.CATEGORY_ID = C.CATEGORY_ID
JOIN SUPPLIERS AS S ON P.SUPPLIER_ID = S.SUPPLIER_ID
WHERE UNIT_PRICE > 50;
```

| PRODUCT_NAME            | CATEGORY_NAME  | COMPANY_NAME                      | UNIT_PRICE |
|-------------------------|----------------|-----------------------------------|------------|
| Chateau de Blaye        | Beverages      | Aux joyeux ecclésiastiques        | 263.5000   |
| Sir Rodney's Marmalade  | Confections    | Specialty Biscuits, Ltd.          | 81.0000    |
| Raclette Courdavault    | Dairy Products | Gai pâturage                      | 55.0000    |
| Mishi Kobe Niku         | Meat/Poultry   | Tokyo Traders                     | 97.0000    |
| Thüringer Rostbratwurst | Meat/Poultry   | Plutzer Lebensmittelgroßmärkte AG | 123.7900   |
| Manjimup Dried Apples   | Produce        | G'day, Mate                       | 53.0000    |
| Carnarvon Tigers        | Seafood        | Pavlova, Ltd.                     | 62.5000    |

# Outer Join

- Consider the tables EMPS and DEPTS

EMPS

| EMPNO | ENAME  | DEPTNO |
|-------|--------|--------|
| 1     | SCOTT  | 1      |
| 2     | MILLER | 2      |
| 3     | ALLEN  | 2      |
| 4     | ADAMS  | 4      |
| 5     | SMITH  | 4      |

DEPTS

| DEPTNO | DNAME      |
|--------|------------|
| 1      | ADMIN      |
| 2      | ACCOUNTING |
| 3      | MARKETING  |
| 4      | OPERATIONS |
| 5      | SALES      |

- As you may have observed, there are no employees in the departments 'MARKETING' and 'SALES'
- The join statement does not get all the departments, since the join condition can not find match for DEPTNO 4 and 5

```
SELECT ENAME, DNAME
FROM EMPS JOIN DEPTS
ON EMPS.DEPTNO = DEPTS.DEPTNO;
```

| ENAME  | DNAME      |
|--------|------------|
| SCOTT  | ADMIN      |
| MILLER | ACCOUNTING |
| ALLEN  | ACCOUNTING |
| ADAMS  | OPERATIONS |
| SMITH  | OPERATIONS |

# Outer Join

- This default behavior is called INNER join

```
SELECT ENAME, DNAME
FROM EMPS JOIN DEPTS
ON EMPS.DEPTNO = DEPTS.DEPTNO;
```

Same As

```
SELECT ENAME, DNAME
FROM EMPS INNER JOIN DEPTS
ON EMPS.DEPTNO = DEPTS.DEPTNO;
```

- If you want those rows from either of the tables that do not have a corresponding row in the other table, then use OUTER JOIN.

```
SELECT ENAME, DNAME
FROM EMPS RIGHT OUTER JOIN DEPTS
ON EMPS.DEPTNO = DEPTS.DEPTNO;
```

Note that DEPTS table, which has excess rows, appears on the RIGHTSIDE of the JOIN keyword

or

```
SELECT ENAME, DNAME
FROM DEPTS LEFT OUTER JOIN EMPS
ON DEPTS.DEPTNO = EMPS.DEPTNO;
```

Note that DEPTS table, which has excess rows, appears on the LEFTSIDE of the JOIN keyword

## Result of OUTER join

| ENAME  | DNAME      |
|--------|------------|
| SCOTT  | ADMIN      |
| MILLER | ACCOUNTING |
| ALLEN  | ACCOUNTING |
| NULL   | MARKETING  |
| ADAMS  | OPERATIONS |
| SMITH  | OPERATIONS |
| NULL   | SALES      |

- You can replace the NULL values using IFNULL() function

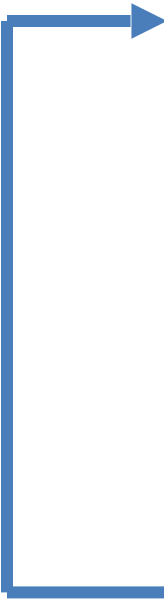
```
SELECT IFNULL(ENAME, '--N/A--') AS ENAME, DNAME
FROM EMPS RIGHT OUTER JOIN DEPTS
ON EMPS.DEPTNO = DEPTS.DEPTNO;
```

| ENAME   | DNAME      |
|---------|------------|
| SCOTT   | ADMIN      |
| MILLER  | ACCOUNTING |
| ALLEN   | ACCOUNTING |
| --N/A-- | MARKETING  |
| ADAMS   | OPERATIONS |
| SMITH   | OPERATIONS |
| --N/A-- | SALES      |

# Self Join

- If a table has a self referencing foreign key, then self join can be used to get relevant data

```
mysql> DESC EMPLOYEES;
```



| Field             | Type        | Null | Key | Default | Extra |
|-------------------|-------------|------|-----|---------|-------|
| EMPLOYEE_ID       | int(11)     | NO   | PRI | NULL    |       |
| LAST_NAME         | varchar(10) | NO   |     | NULL    |       |
| FIRST_NAME        | varchar(10) | NO   |     | NULL    |       |
| TITLE             | varchar(25) | YES  |     | NULL    |       |
| TITLE_OF_COURTESY | varchar(5)  | YES  |     | NULL    |       |
| BIRTH_DATE        | datetime    | YES  |     | NULL    |       |
| HIRE_DATE         | datetime    | YES  |     | NULL    |       |
| ADDRESS           | varchar(30) | YES  |     | NULL    |       |
| CITY              | varchar(10) | YES  |     | NULL    |       |
| REGION            | varchar(2)  | YES  |     | NULL    |       |
| POSTAL_CODE       | varchar(10) | YES  |     | NULL    |       |
| COUNTRY           | varchar(10) | YES  |     | NULL    |       |
| HOME_PHONE        | varchar(15) | YES  |     | NULL    |       |
| EXTENSION         | varchar(5)  | YES  |     | NULL    |       |
| Photo             | longblob    | YES  |     | NULL    |       |
| NOTES             | longtext    | YES  |     | NULL    |       |
| REPORTS_TO        | int(11)     | YES  |     | NULL    |       |

# Self Join

EMPLOYEES (AS EMPS)

| EMPLOYEE_ID | FIRST_NAME | LAST_NAME | REPORTS_TO |
|-------------|------------|-----------|------------|
| 1           | Nancy      | Davolio   | 2          |
| 2           | Andrew     | Fuller    | NULL       |
| 3           | Janet      | Leverling | 2          |
| 4           | Margaret   | Peacock   | 2          |
| 5           | Steven     | Buchanan  | 2          |
| 6           | Michael    | Suyama    | 5          |
| 7           | Robert     | King      | 5          |
| 8           | Laura      | Callahan  | 2          |
| 9           | Anne       | Dodsworth | 5          |

EMPLOYEES (AS MANAGERS)

| EMPLOYEE_ID | FIRST_NAME | LAST_NAME |
|-------------|------------|-----------|
| 1           | Nancy      | Davolio   |
| 2           | Andrew     | Fuller    |
| 3           | Janet      | Leverling |
| 4           | Margaret   | Peacock   |
| 5           | Steven     | Buchanan  |
| 6           | Michael    | Suyama    |
| 7           | Robert     | King      |
| 8           | Laura      | Callahan  |
| 9           | Anne       | Dodsworth |

- From the above data, we can observe that Nancy(1), Janet(3), Margaret(4), Steven(5), and Laura(8) report to Andrew(2). Also, Michael(6), Robert(7), and Anne(9) report to Steven(5).
-

# Self Join

EMPLOYEES (AS EMPS)

| EMPLOYEE_ID | FIRST_NAME | LAST_NAME | REPORTS_TO |
|-------------|------------|-----------|------------|
| 1           | Nancy      | Davolio   | 2          |
| 2           | Andrew     | Fuller    | NULL       |
| 3           | Janet      | Leverling | 2          |
| 4           | Margaret   | Peacock   | 2          |
| 5           | Steven     | Buchanan  | 2          |
| 6           | Michael    | Suyama    | 5          |
| 7           | Robert     | King      | 5          |
| 8           | Laura      | Callahan  | 2          |
| 9           | Anne       | Dodsworth | 5          |



EMPLOYEES (AS MANAGERS)

| EMPLOYEE_ID | FIRST_NAME | LAST_NAME |
|-------------|------------|-----------|
| 1           | Nancy      | Davolio   |
| 2           | Andrew     | Fuller    |
| 3           | Janet      | Leverling |
| 4           | Margaret   | Peacock   |
| 5           | Steven     | Buchanan  |
| 6           | Michael    | Suyama    |
| 7           | Robert     | King      |
| 8           | Laura      | Callahan  |
| 9           | Anne       | Dodsworth |



| FIRST_NAME<br>(EMPLOYEE)                      | REPORTS_TO / EMPLOYEE_ID | FIRST_NAME<br>(MANAGER) |
|-----------------------------------------------|--------------------------|-------------------------|
| Nancy<br>Janet<br>Margaret<br>Steven<br>Laura | 2                        | Andrew                  |
| Michael<br>Robert<br>Anne                     | 5                        | Steven                  |

```
SELECT EMPS.FIRST_NAME, MANAGERS.FIRST_NAME
FROM EMPLOYEES AS EMPS JOIN EMPLOYEES AS MANAGERS
ON EMPS.REPORTS_TO = MANAGERS.EMPLOYEE_ID;
```



# Self Join - Examples

- Retrieve the PRODUCT\_NAME, UNIT\_PRICE of all the products that are costlier than the product with PRODUCT\_ID 62

```
SELECT P1.PRODUCT_NAME, P1.UNIT_PRICE
FROM PRODUCTS P1 JOIN PRODUCTS P2
ON P1.UNIT_PRICE > P2.UNIT_PRICE
WHERE P2.PRODUCT_ID = 62;
```

| PRODUCT_NAME            | UNIT_PRICE |
|-------------------------|------------|
| Mishi Kobe Niku         | 97.0000    |
| Carnarvon Tigers        | 62.5000    |
| Sir Rodney's Marmalade  | 81.0000    |
| Thüringer Rostbratwurst | 123.7900   |
| Chte de Blaye           | 263.5000   |
| Manjimup Dried Apples   | 53.0000    |
| Raclette Courdavault    | 55.0000    |

# Self Join - Examples

- Retrieve the EMPLOYEE\_ID, FIRST\_NAME, LAST\_NAME, TITLE of all employees who are from the same city as the employee with FIRST\_NAME 'STEVEN'

```
SELECT E1.EMPLOYEE_ID, E1.FIRST_NAME, E1.LAST_NAME, E1.TITLE
FROM EMPLOYEES E1 JOIN EMPLOYEES E2
ON E1.CITY = E2.CITY
WHERE E2.FIRST_NAME = 'STEVEN';
```

| EMPLOYEE_ID | FIRST_NAME | LAST_NAME | TITLE                |
|-------------|------------|-----------|----------------------|
| 5           | Steven     | Buchanan  | Sales Manager        |
| 6           | Michael    | Suyama    | Sales Representative |
| 7           | Robert     | King      | Sales Representative |
| 9           | Anne       | Dodsworth | Sales Representative |

Using

# Subqueries

---

# Using a subquery to solve a problem

- Which are the products costlier than average price of all products?

Main query:

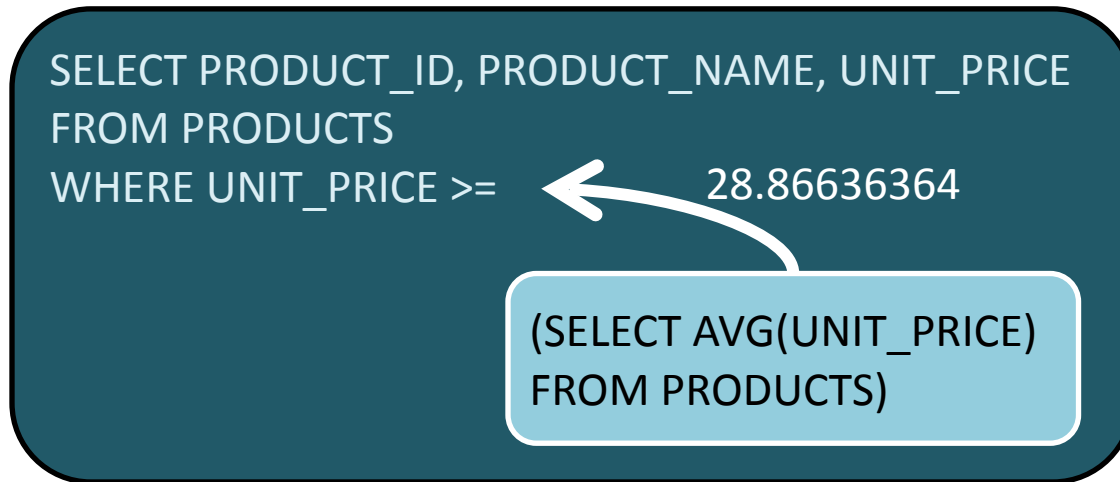
Which are the products whose unit\_price is  
More than average price of all products?

Subquery:

What is the average price of all products?

---

# Using a subquery to solve a problem



- A subquery is a SELECT statement embedded in the WHERE clause of another SELECT statement
  - The subquery is executed once before the main query is executed
  - The result of the subquery is substituted for the subquery and is used by the main query for every row
-

# Guidelines for Using Subqueries

- Enclose subqueries in parentheses.
  - Place subqueries on the right side of the comparison condition.
  - The ORDER BY clause in the subquery is not needed unless you are performing Top-N analysis
  - Use single-row operators with single-row subqueries and use multiple-row operators with multiple-row subqueries.
-

# Multi-row subqueries

- If we have to retrieve all the products that belong to those category of products supplied by a supplier with SUPPLIER\_ID 6

- Subquery:

Get the category ids from products table where supplier id is 6:

```
SELECT DISTINCT CATEGORY_ID
FROM PRODUCTS
WHERE SUPPLIER_ID=6;
```

- This may result in the category ids – 2, 7, 8
- This type of subqueries are known as multi-row subqueries, and an outer query (or the main query) must use 'IN' or 'NOT IN' to compare them

```
SELECT PRODUCT_ID, PRODUCT_NAME, CATEGORY_ID
FROM PRODUCTS WHERE CATEGORY_ID IN
(SELECT DISTINCT CATEGORY_ID FROM PRODUCTS WHERE SUPPLIER_ID=6)
ORDER BY CATEGORY_ID;
```

---

# Multi-row subqueries

- Using the equals operator to do the same would result in an error
  - One CATEGORY\_ID from each row can not be compared with a bunch of CATEGORY\_ID values given by the subquery

```
mysql> SELECT PRODUCT_ID, PRODUCT_NAME, CATEGORY_ID
-> FROM PRODUCTS WHERE CATEGORY_ID =
-> (SELECT DISTINCT CATEGORY_ID FROM PRODUCTS WHERE SUPPLIER_ID=6)
-> ORDER BY CATEGORY_ID;
ERROR 1242 (21000): Subquery returns more than 1 row
```

- Alternatively, you can use  
= ANY  
operator to do the same

```
SELECT PRODUCT_ID, PRODUCT_NAME, CATEGORY_ID
FROM PRODUCTS WHERE CATEGORY_ID = ANY
(SELECT DISTINCT CATEGORY_ID FROM PRODUCTS WHERE SUPPLIER_ID=6)
ORDER BY CATEGORY_ID;
```



# Multi-Column Subqueries

- Get the details of the products that belong to the same category and are supplied by the same supplier as that of the product with id '6'

PRODUCT\_ID = 6

| SUPPLIER_ID | CATEGORY_ID |
|-------------|-------------|
| 3           | 2           |

- The outer/main query can use two columns together to compare the result of the subquery

```
SELECT PRODUCT_ID, PRODUCT_NAME, CATEGORY_ID, SUPPLIER_ID
FROM PRODUCTS WHERE (SUPPLIER_ID, CATEGORY_ID) =
(SELECT SUPPLIER_ID, CATEGORY_ID FROM PRODUCTS WHERE PRODUCT_ID = 6);
```

| PRODUCT_ID | PRODUCT_NAME                 | CATEGORY_ID | SUPPLIER_ID |
|------------|------------------------------|-------------|-------------|
| 6          | Grandma's Boysenberry Spread | 2           | 3           |
| 8          | Northwoods Cranberry Sauce   | 2           | 3           |

# Multi-Column Subqueries

```
SELECT PRODUCT_ID, PRODUCT_NAME, CATEGORY_ID, SUPPLIER_ID
FROM PRODUCTS WHERE (SUPPLIER_ID, CATEGORY_ID) =
(SELECT SUPPLIER_ID, CATEGORY_ID FROM PRODUCTS WHERE PRODUCT_ID = 6);
```



WHERE (SUPPLIER\_ID, CATEGORY\_ID) = (3, 2)

The diagram illustrates the execution of the multi-column subquery. It shows a light blue box containing the text 'WHERE (SUPPLIER\_ID, CATEGORY\_ID) = (3, 2)'. Two blue arrows originate from this box: one starts at the first value '3' and points to the 'SUPPLIER\_ID' column in the main query's WHERE clause, and the other starts at the second value '2' and points to the 'CATEGORY\_ID' column. This visualizes how the subquery's results are used to filter the main query's data.

# Using a Subquery in the FROM Clause

- Retrieve the CATEGORY\_NAME, DESCRIPTION, NUMBER\_OF\_PRODUCTS, AVG\_PRICE\_OF\_PRODUCT for each category in the CATEGORIES table

```
SELECT CATEGORY_NAME, DESCRIPTION, T1.NUMBER_OF_PRODUCTS, T1.AVG_PRICE
FROM CATEGORIES C1 JOIN
 (SELECT CATEGORY_ID, COUNT(*) NUMBER_OF_PRODUCTS,
 AVG(UNIT_PRICE) AVG_PRICE
 FROM PRODUCTS GROUP BY CATEGORY_ID) AS T1
ON C1.CATEGORY_ID = T1.CATEGORY_ID;
```

| CATEGORY_NAME  | DESCRIPTION                                                | NUMBER_OF_PRODUCTS | AVG_PRICE   |
|----------------|------------------------------------------------------------|--------------------|-------------|
| Beverages      | Soft drinks, coffees, teas, beers, and ales                | 12                 | 37.97915667 |
| Condiments     | Sweet and savory sauces, relishes, spreads, and seasonings | 12                 | 23.06250000 |
| Confections    | Desserts, candies, and sweet breads                        | 13                 | 25.16000000 |
| Dairy Products | Cheeses                                                    | 10                 | 28.73000000 |
| Grains/Cereals | Breads, crackers, pasta, and cereal                        | 7                  | 20.25000000 |
| Meat/Poultry   | Prepared meats                                             | 6                  | 54.00565667 |
| Produce        | Dried fruit and bean curd                                  | 5                  | 32.37000000 |
| Seafood        | Seaweed and fish                                           | 12                 | 20.68250000 |

Note: This is only for illustration. The same can be done using a simple joining of CATEGORIES table with PRODUCTS table.

# Correlated Subqueries

- A correlated subquery is a subquery that contains a reference to a table that also appears in the outer query

```
SELECT * FROM t1 WHERE column1 =
 (SELECT column1 FROM t2 WHERE t2.column2 = t1.column2);
```

- Notice that the subquery contains a reference to a column of t1, even though the subquery's FROM clause does not mention a table t1.
  - So, MySQL looks outside the subquery, and finds t1 in the outer query.
  - MySQL evaluates the inner query for each of the row in the outer query, unlike the rest of the cases, where the inner query is evaluated only once for all the rows of the outer query.
-

# Correlated Subqueries

- For certain cases, a correlated subquery is optimized
  - Otherwise, they are inefficient and likely to be slow.
  - Rewriting the query as a join might improve performance.
-

# Correlated Subqueries - Example

- Retrieve the costliest product in each category. Display these fields:  
Category name, Product Name, Unit price

```
SELECT CATEGORY_NAME, PRODUCT_NAME, UNIT_PRICE
FROM CATEGORIES C JOIN PRODUCTS P
ON C.CATEGORY_ID = P.CATEGORY_ID
WHERE UNIT_PRICE =
 (SELECT MAX(UNIT_PRICE) FROM PRODUCTS
 WHERE CATEGORY_ID = P.CATEGORY_ID)
ORDER BY CATEGORY_NAME;
```

| CATEGORY_NAME  | PRODUCT_NAME            | UNIT_PRICE |
|----------------|-------------------------|------------|
| Beverages      | Chate de Blaye          | 263.5000   |
| Condiments     | Vegie-spread            | 43.9000    |
| Confections    | Sir Rodney's Marmalade  | 81.0000    |
| Dairy Products | Raclette Courdavault    | 55.0000    |
| Grains/Cereals | Gnocchi di nonna Alice  | 38.0000    |
| Meat/Poultry   | Thuringer Rostbratwurst | 123.7900   |
| Produce        | Manjimup Dried Apples   | 53.0000    |
| Seafood        | Carnarvon Tigers        | 62.5000    |

# Views

---

# Database View

- A view consists of a stored query accessible as a virtual table in a relational database.
  - Unlike ordinary tables (base tables) in a relational database, a view does not form part of the physical schema:
    - It is a dynamic, virtual table computed or collated from data in the database.
  - Changing the data in a table alters the data shown in subsequent invocations of the view.
-



# Advantages of a view (1 of 2)

- Views can represent a subset of the data contained in a table
  - Views can join and simplify multiple tables into a single virtual table
  - Views can act as aggregated tables, where the database engine aggregates data (sum, average etc.) and presents the calculated results as part of the data
  - Views can hide the complexity of data; for example a view could appear as Sales2000 or Sales2001, transparently partitioning the actual underlying table
-

## Advantages of a view (2 of 2)

- Views take very little space to store; the database contains only the definition of a view, not a copy of all the data it presents
  - Depending on the SQL engine used, views can provide extra security
  - Views can limit the degree of exposure of a table or tables to the outer world
-

# Creating a view in MySQL

Syntax:

```
CREATE [OR REPLACE]
VIEW view_name [(column_list)]
AS select_statement
```

Example:

```
CREATE VIEW order_amounts_view AS
SELECT o.order_id,
format(sum(od.unit_price*od.quantity*(1-od.discount)),2)
AS order_amount
FROM orders as o JOIN order_details AS od
ON o.order_id = od.order_id
GROUP BY o.order_id;
```

---

## Output from a view

```
mysql> select * from order_amounts_view limit 10;
```

| order_id | order_amount |
|----------|--------------|
| 10248    | 440.00       |
| 10250    | 1,912.60     |
| 10251    | 654.06       |
| 10252    | 3,597.90     |
| 10253    | 1,444.80     |
| 10254    | 556.62       |
| 10255    | 2,490.50     |
| 10256    | 517.80       |
| 10258    | 1,614.88     |
| 10260    | 1,504.65     |

```
10 rows in set (0.08 sec)
```

# Database Index

---

# Database index

- A database index is a data structure that improves the speed of data retrieval operations on a database table at the cost of slower writes and increased storage space.
  - Indices can be created using one or more columns of a database table, providing the basis for both rapid random lookups and efficient access of ordered records.
  - Without an index, MySQL must begin with the first row and then read through the entire table to find the relevant rows.
-

# Database index

Syntax:

```
CREATE INDEX index_name
ON tbl_name (index_col_name,...)
```

Example:

```
CREATE INDEX idx_products_unit_price
ON products (unit_price);
```

To delete the index, use the DROP INDEX command:

```
DROP INDEX idx_products_unit_price ON products;
```

---

# Database index

- In MySQL, the following are automatically indexed:
    - Primary key column
    - Unique column
    - Fulltext column
-



